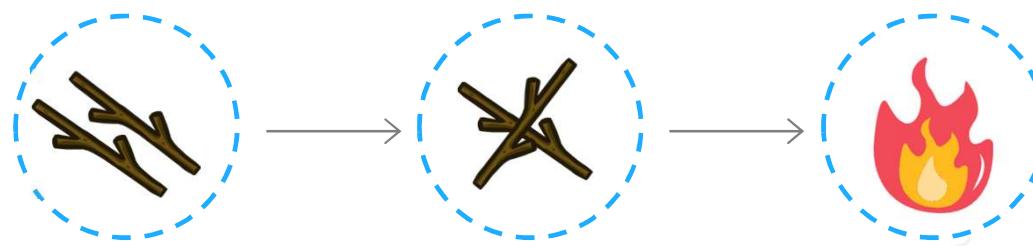
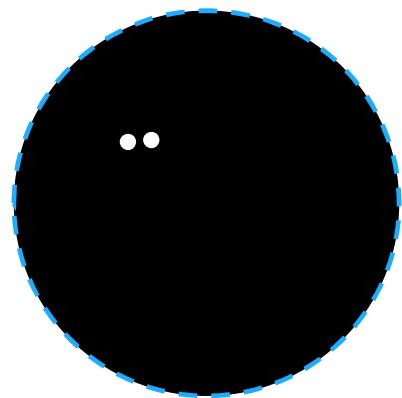




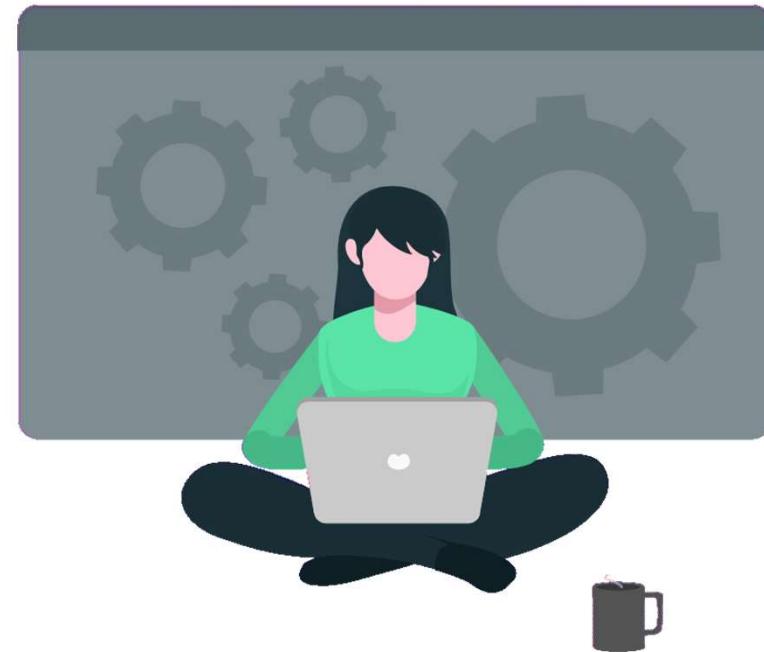
KodeKloud

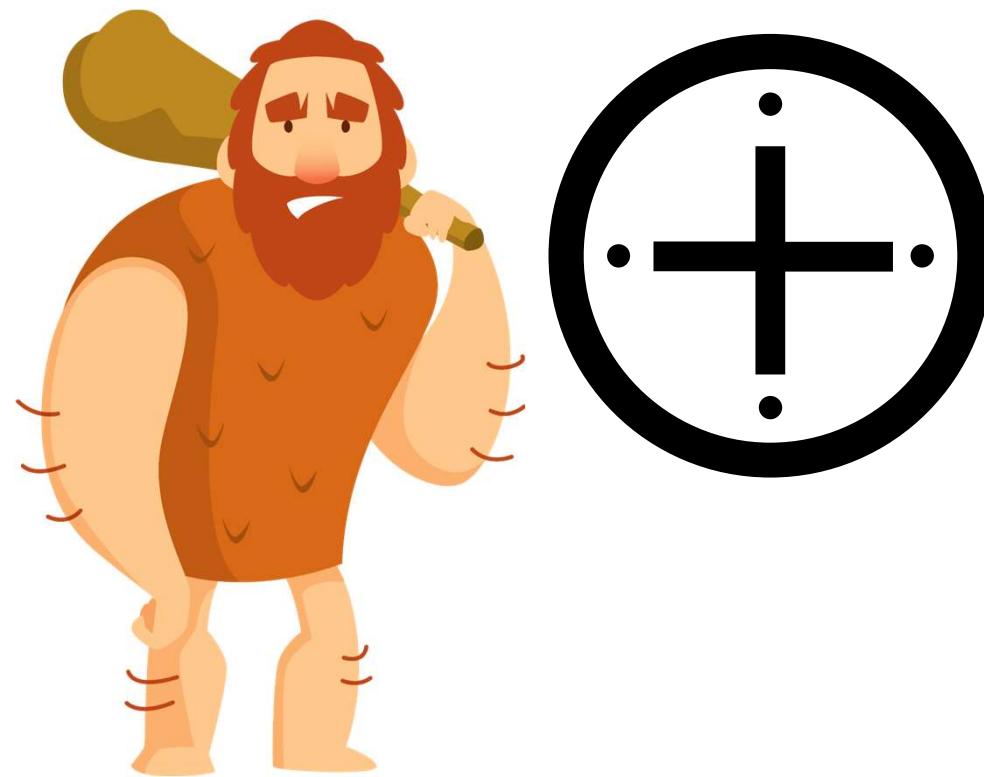
Lambda Service Basics

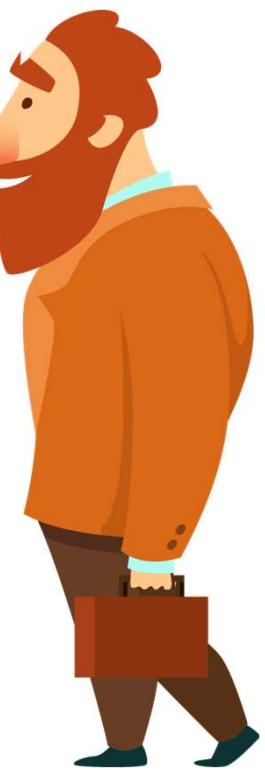


















AWS Lambda

Bring your own Code





What is Lambda?



Serverless
Compute
Service



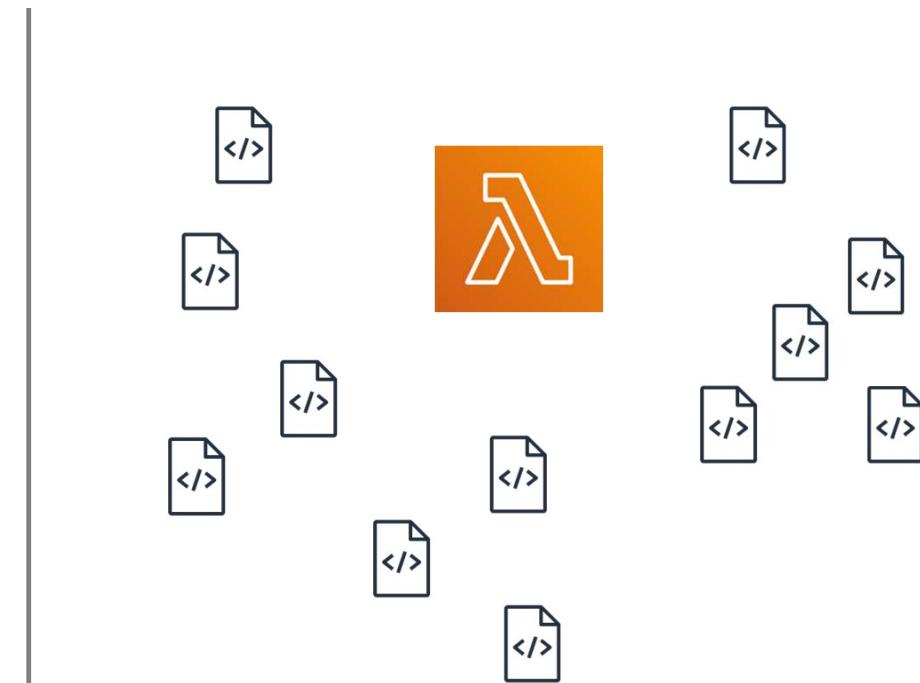


What is Lambda?





What is Lambda?

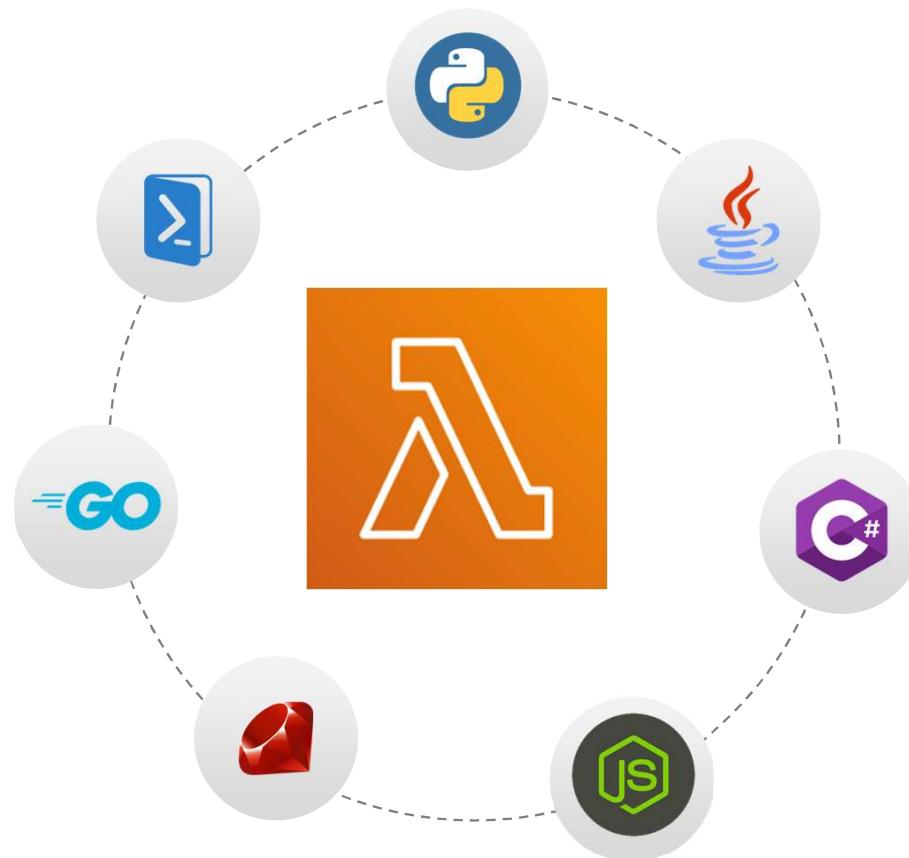




What is Lambda?



Supported Languages



➤ Supported Languages



➤ Supported Languages





Summary



What is AWS Lambda?

AWS Lambda is a serverless computing service that allows you to run programs without having to worry about provisioning, maintaining, or waiting for servers to be built

AWS Lambda Features

- **Serverless:** Run code without provisioning or maintaining a server.
- **Automatic Scaling:** Scale your applications automatically as per the workload.
- **Pay per use:** Billed per millisecond of use.
- **Consistency:** Performance consistency is achieved by selecting the right memory size for the function.
- **Language support:** AWS Lambda supports multiple programming languages.

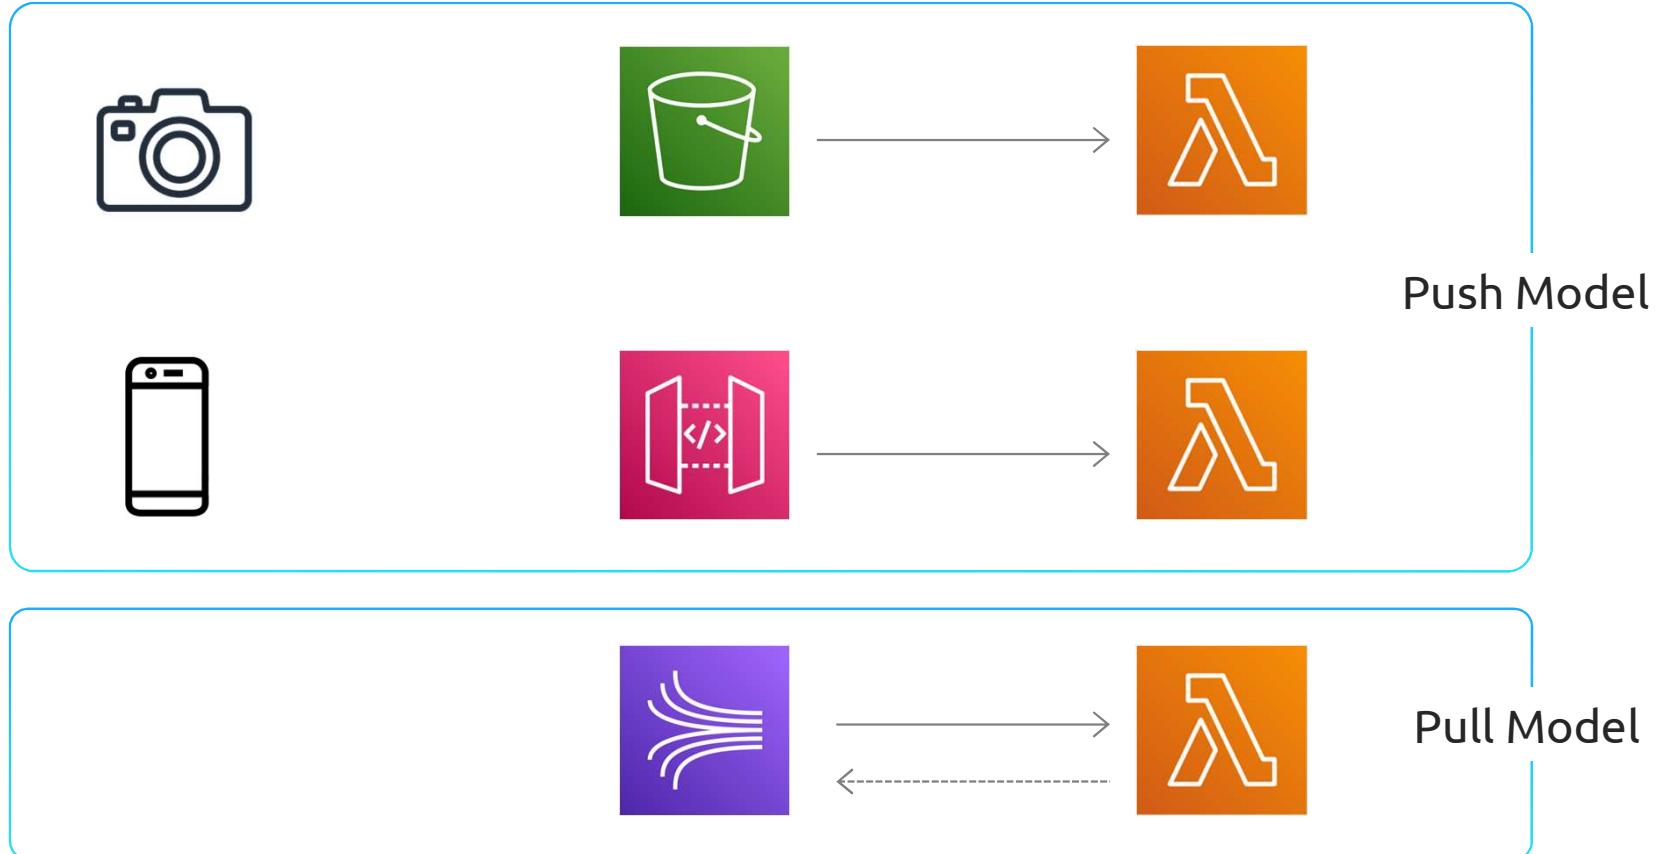


KodeKloud

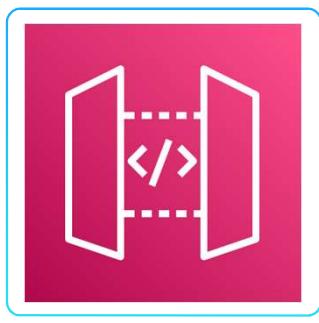
Event Sources



Event Sources



Push Model Source Types : Synchronous



API Gateway

http request

Erroneous http request
Response

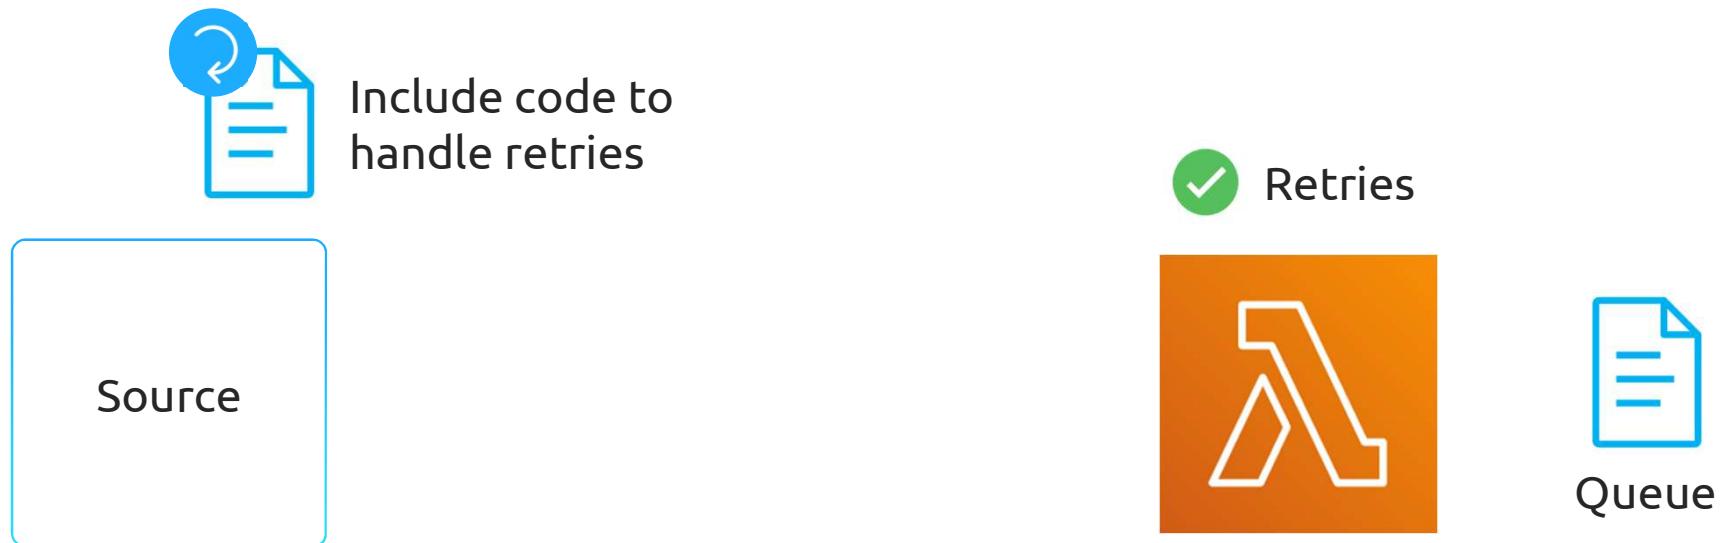
http response

✖ Retries



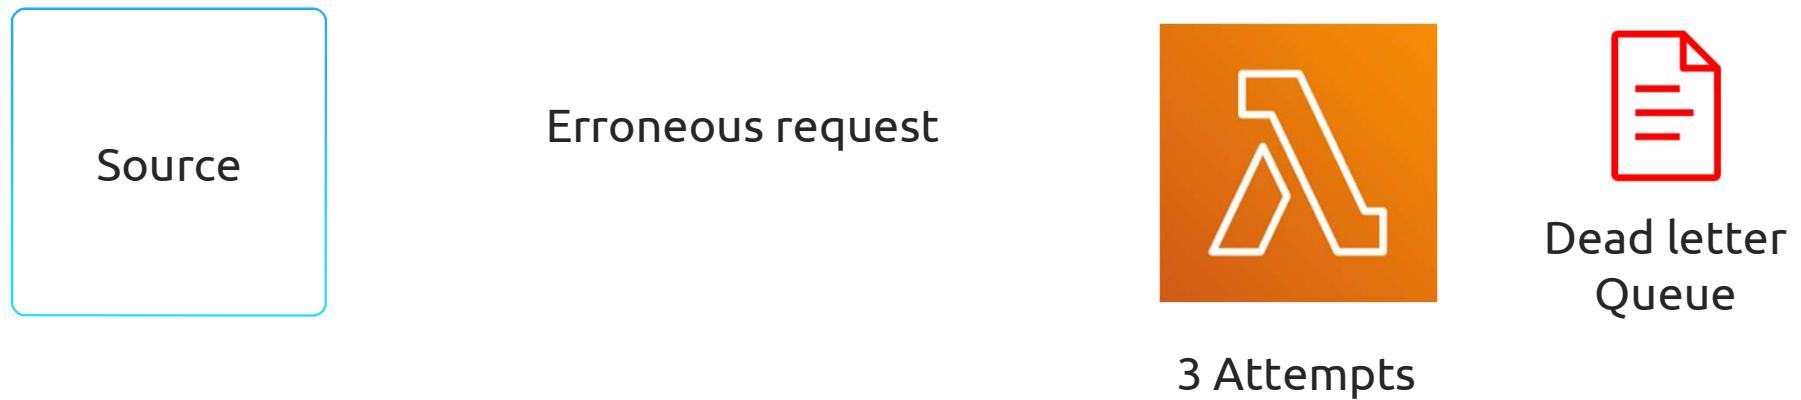


Push Model Source Types : Asynchronous





Push Model Source Types : Asynchronous





Push Model Source Types : Asynchronous



Source



Destinations



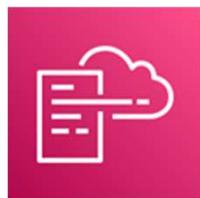
Eventbridge



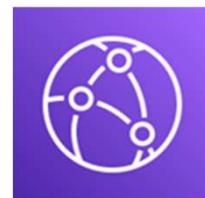
Simple Notification
Service

Push Model Source Types

Synchronous



Cloudformation



Cloudfront



API Gateway



Cognito

Asynchronous



Cloudwatch



Eventbridge



S3



SNS

Pull Model Source Types

Streams



Dynamo DB

Kinesis Data
Stream

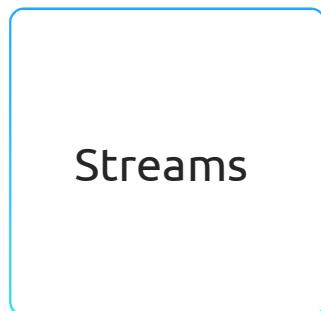
Queues



Amazon Simple
Queueing Service

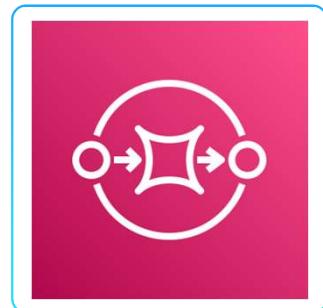


Pull Model Source Types : Streams





Pull Model Source Types : Queues



Amazon Simple
Queueing Service

Queue



Pull Model Source Types



Lambda Function

Pull Model Source Types

Queues



Amazon Simple
Queueing Service

Streams



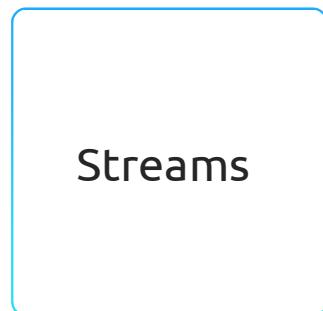
Dynamo DB



Kinesis Data
Stream



Pull Model Source Types : Streams

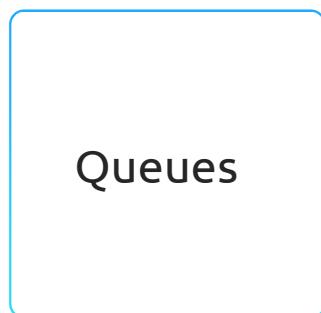


Retry





Pull Model Source Types : Queues



3



Summary

Event Source

AWS Lambda Event Source is any AWS Service, custom application, a stream of data or queue which triggers a Lambda function to run.



Lambda Event Sources Type

- Push-based model
 - Synchronous
 - Asynchronous
- Pull-based model
 - Stream
 - Queue



Summary

- **Push-based model:** Other service directly triggers Lambda and tells it to activate when something happens.
 - Synchronous: Lambda returns a response back to the event source.
 - Asynchronous: For asynchronous invocations, Lambda does handle retries. After Lambda is invoked by the source, it will first place the event into a queue and immediately sends a success response back to the source.
- **Pull-based model:** Information is flowing through a stream or put in a queue that Lambda periodically polls, and it then goes into action when certain events are detected .
 - Stream: Lambda will stop polling while it retries the message.
 - Queue: Lambda will return the message to the queue if the invocation fails or times out but will also keep retrying until it's either successful or expires.

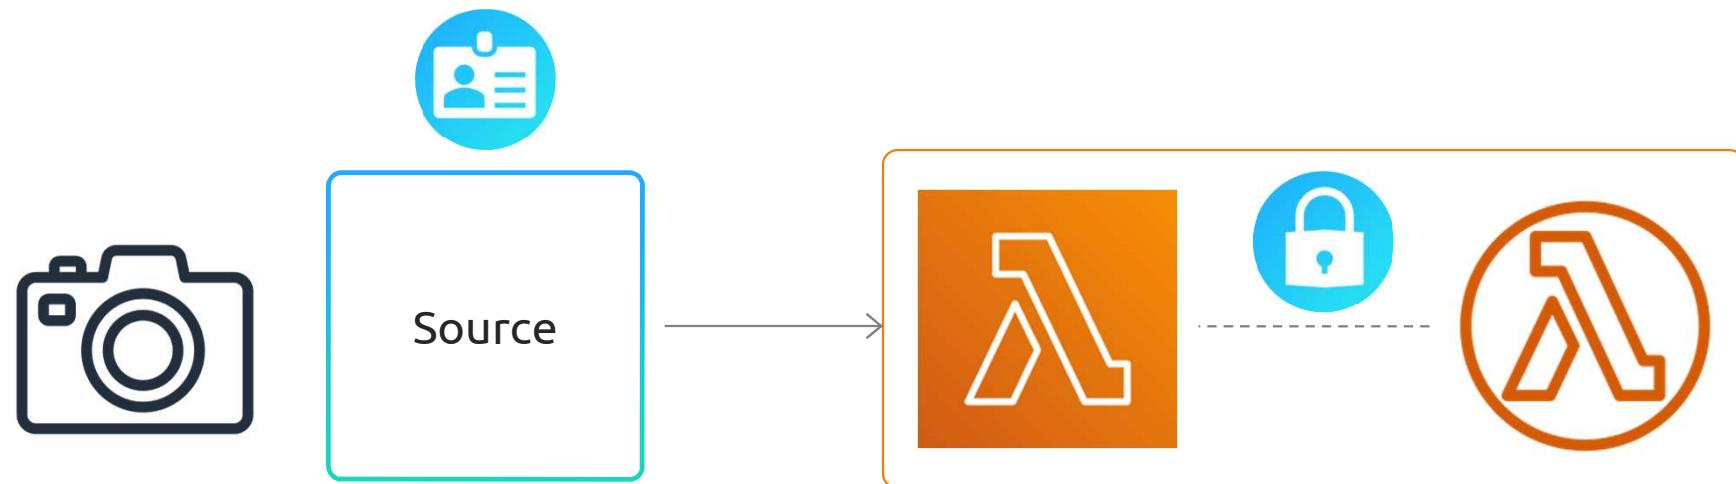


KodeKloud

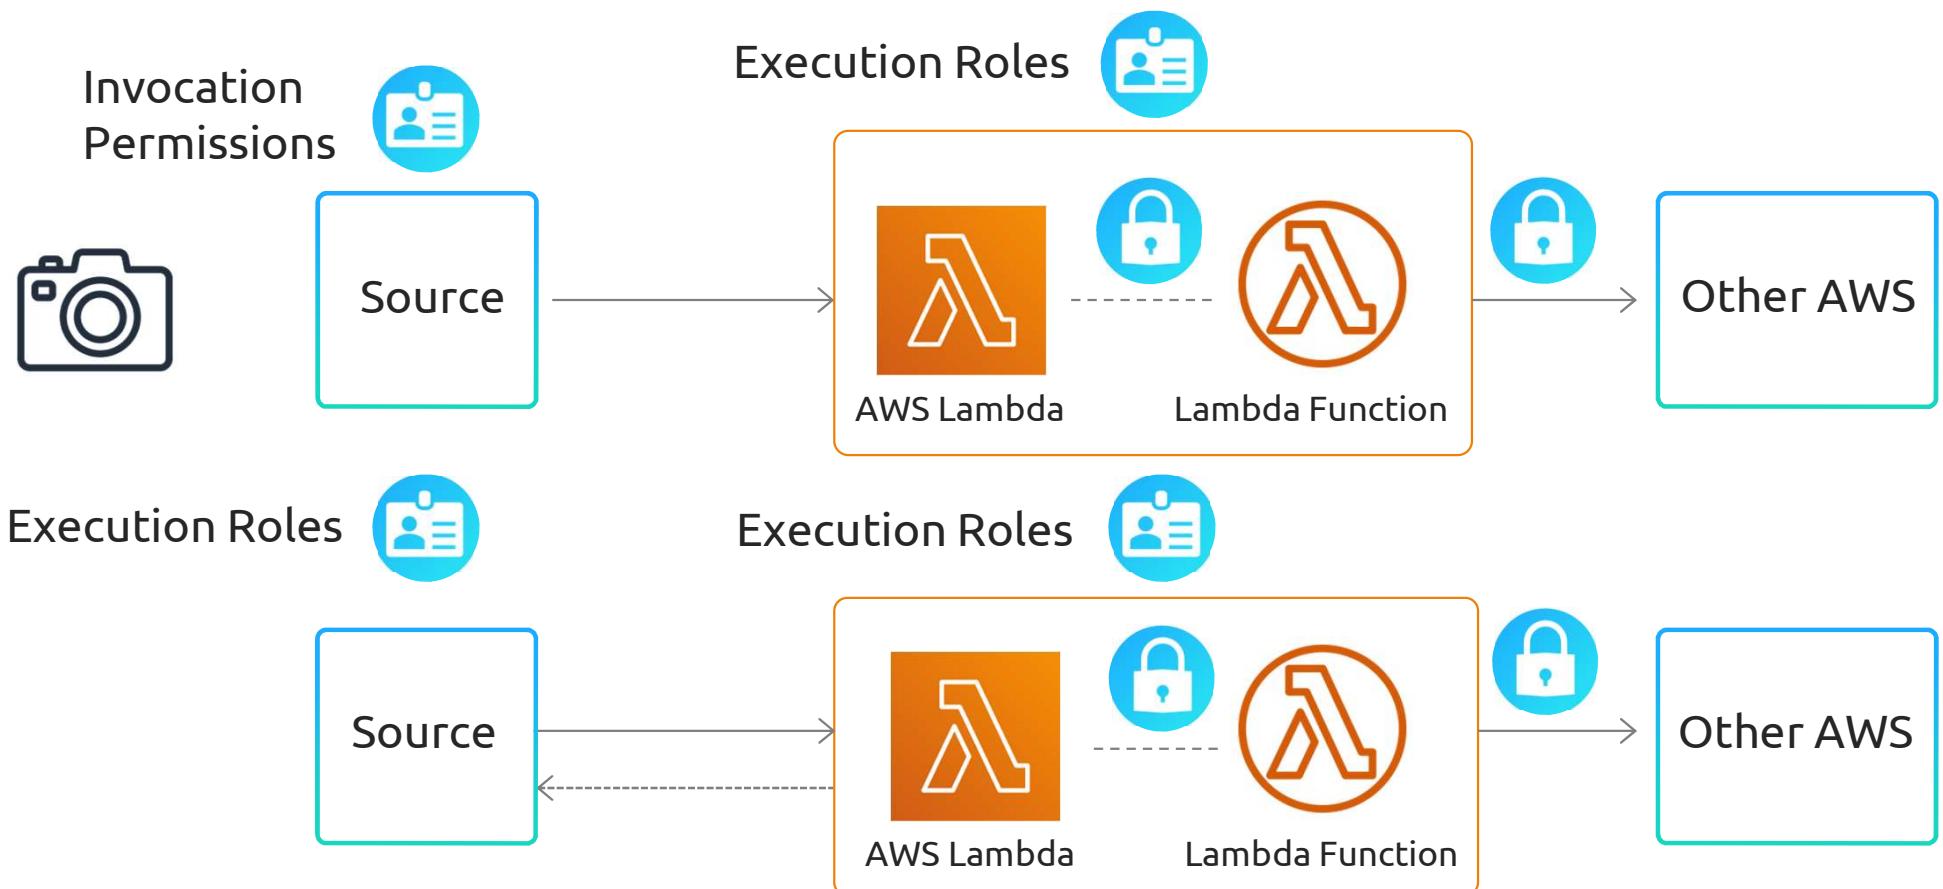
Access Permissions

Access Permissions

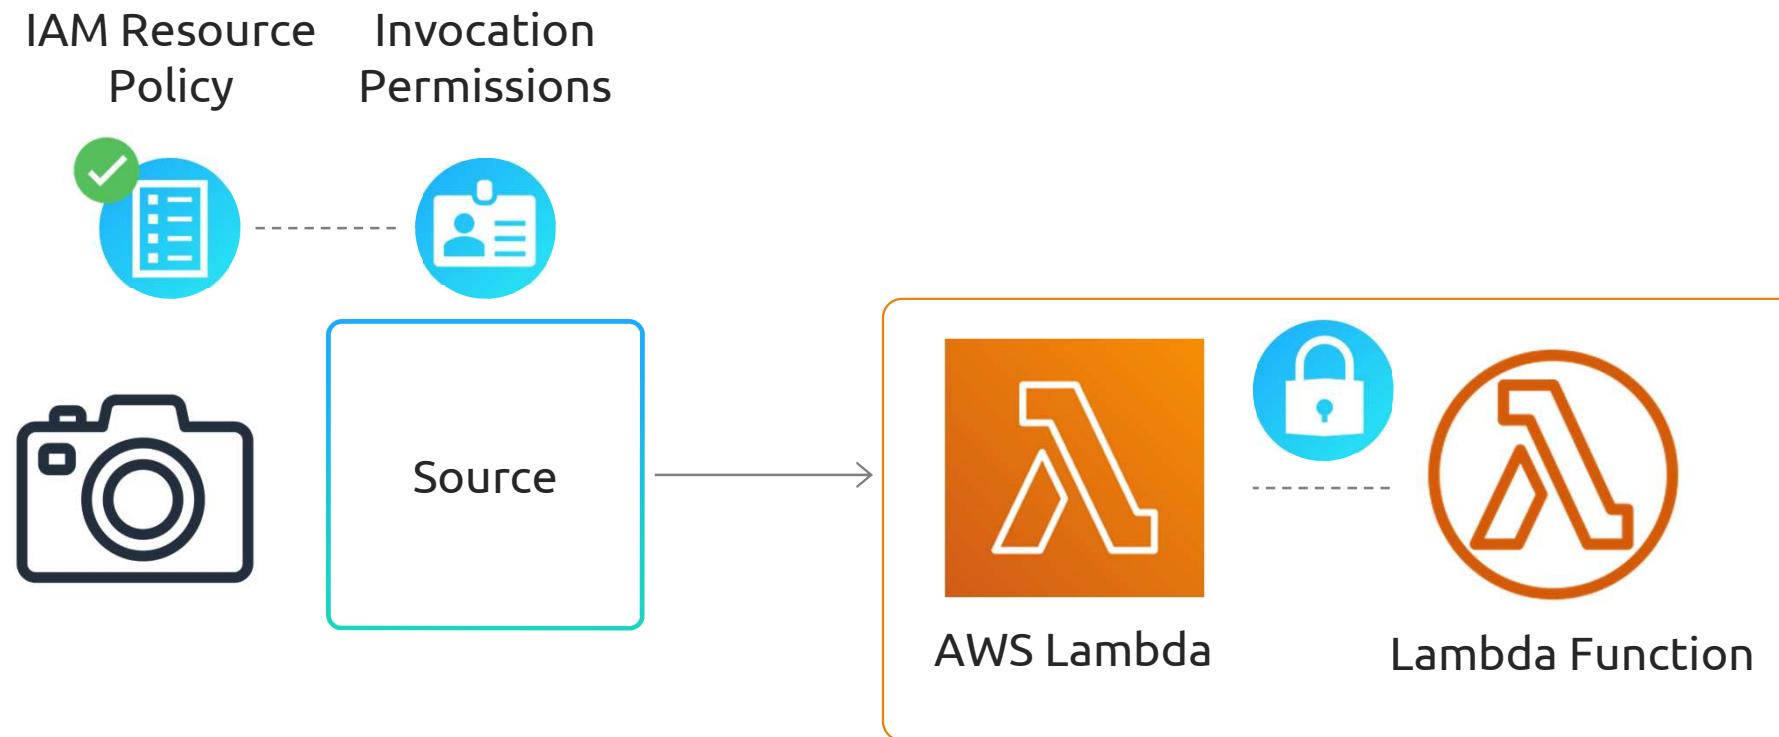
Invocation Permissions



Access Permissions

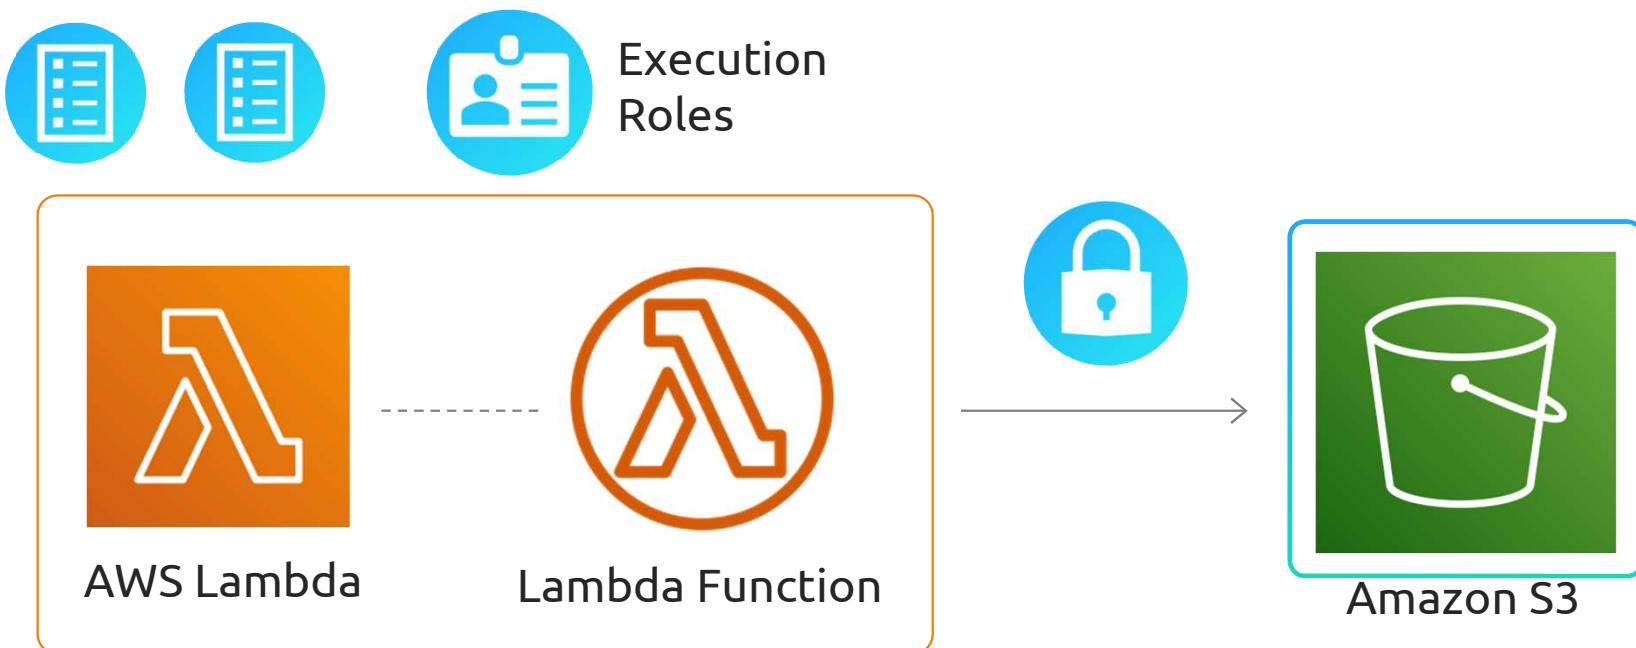


Invocation Permissions



Execution Roles

2 IAM Resource Policies



IAM Policy

```
Terminal

{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "ExampleSourceFunctionArn",
            "Effect": "Allow",
            "Action": "s3:PutObject",
            "Resource": "arn:aws:s3:::lambda_bucket/*",
            "Condition": {
                "ArnEquals": {
                    "lambda:SourceFunctionArn": "arn:aws:lambda:us-east-1:123456789012:function:source_lambda"
                }
            }
        }
    ]
}
```



Trust Policy

```
Terminal

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Service": "lambda.amazonaws.com"
      },
      "Action": "sts:AssumeRole"
    }
  ]
}
```





Summary



Access Permissions

- Security is crucial when dealing with Lambda because of its power to run code and take actions that affect other AWS services.
- Two types of security permissions are needed for Lambda: [Invocation permissions](#) and [Execution roles](#).
- [Invocation permissions](#) are only needed for [push event sources](#) and are [granted through an IAM resource policy](#) automatically created when configuring an AWS service as an event source.
- [Execution roles](#) grant Lambda permissions to [interact with other AWS services](#). They include an [IAM policy](#) defining what Lambda is allowed to do and a [Trust policy](#) allowing Lambda to assume the execution role and perform actions on the other service.
- Adding the execution role to your Lambda function is the final step in granting permissions and policies.

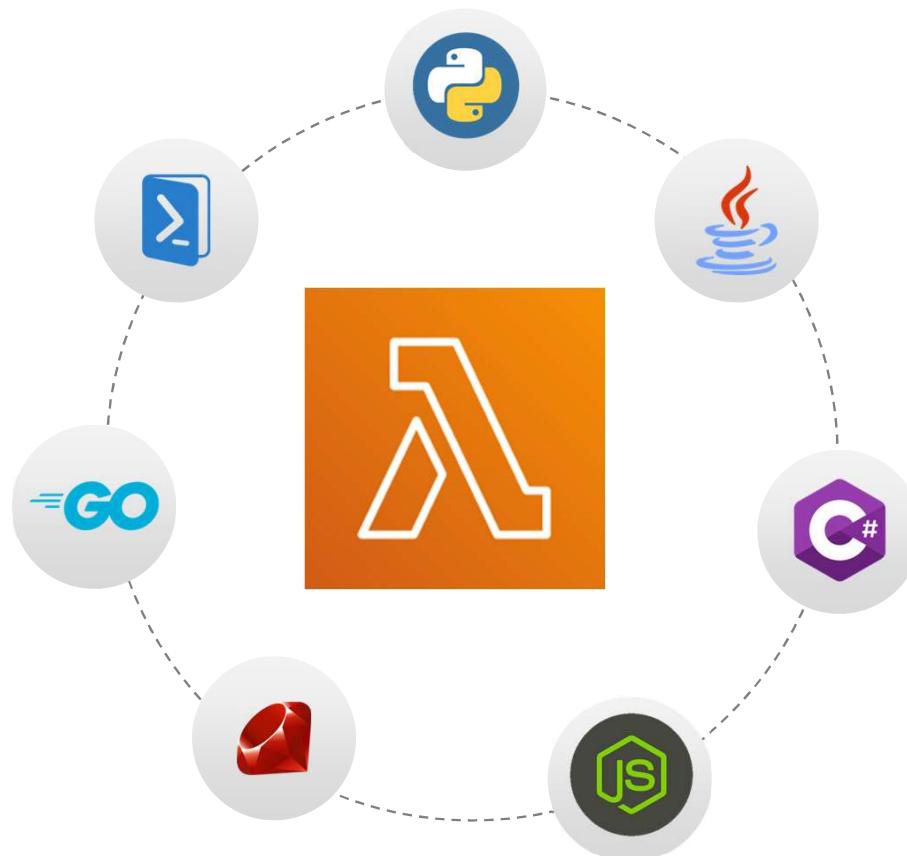


KodeKloud

Functions

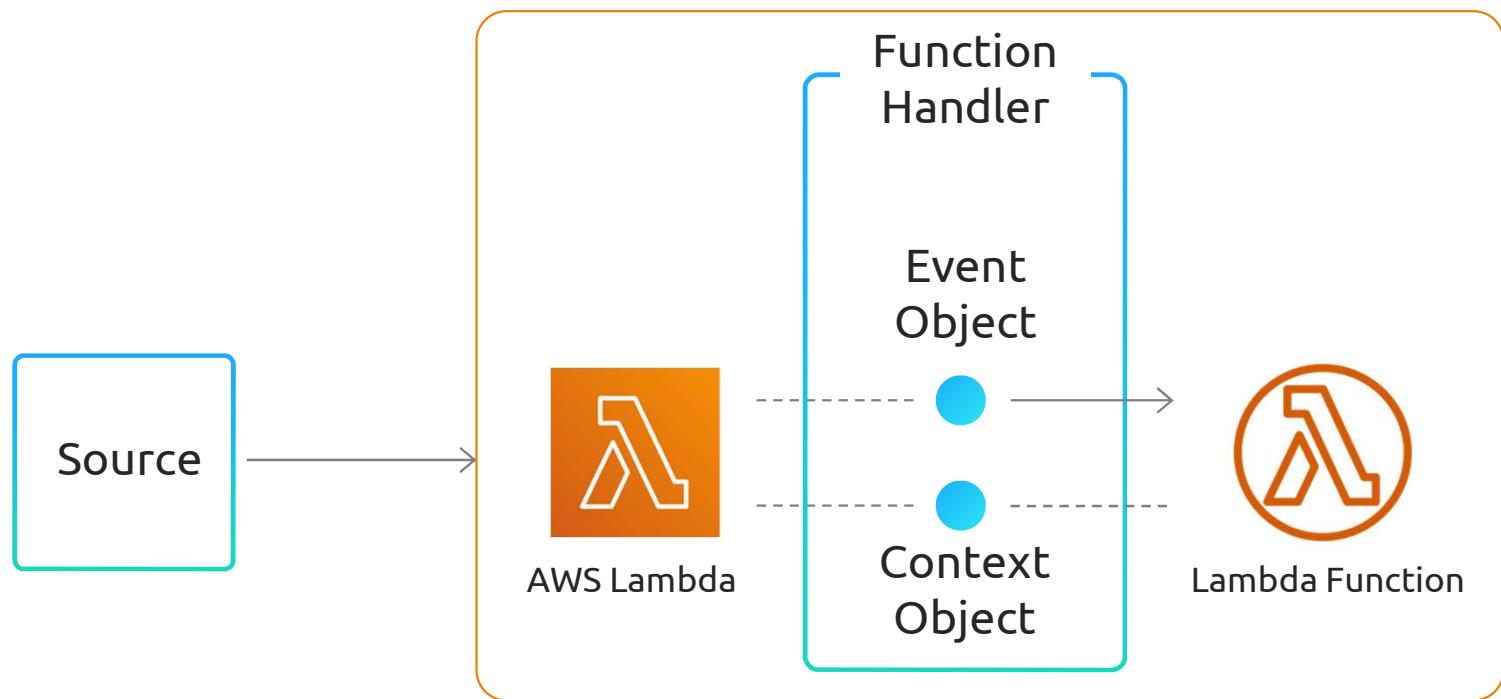


Functions



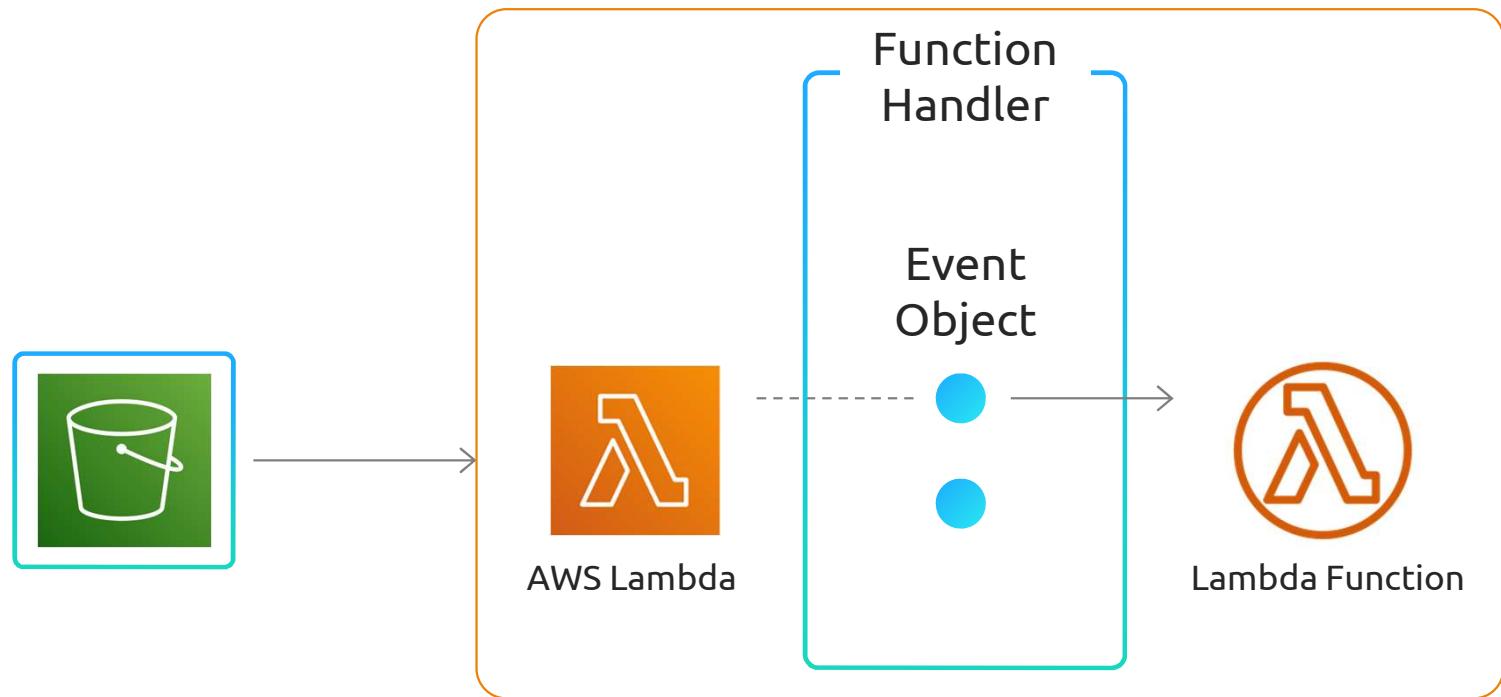


Functions



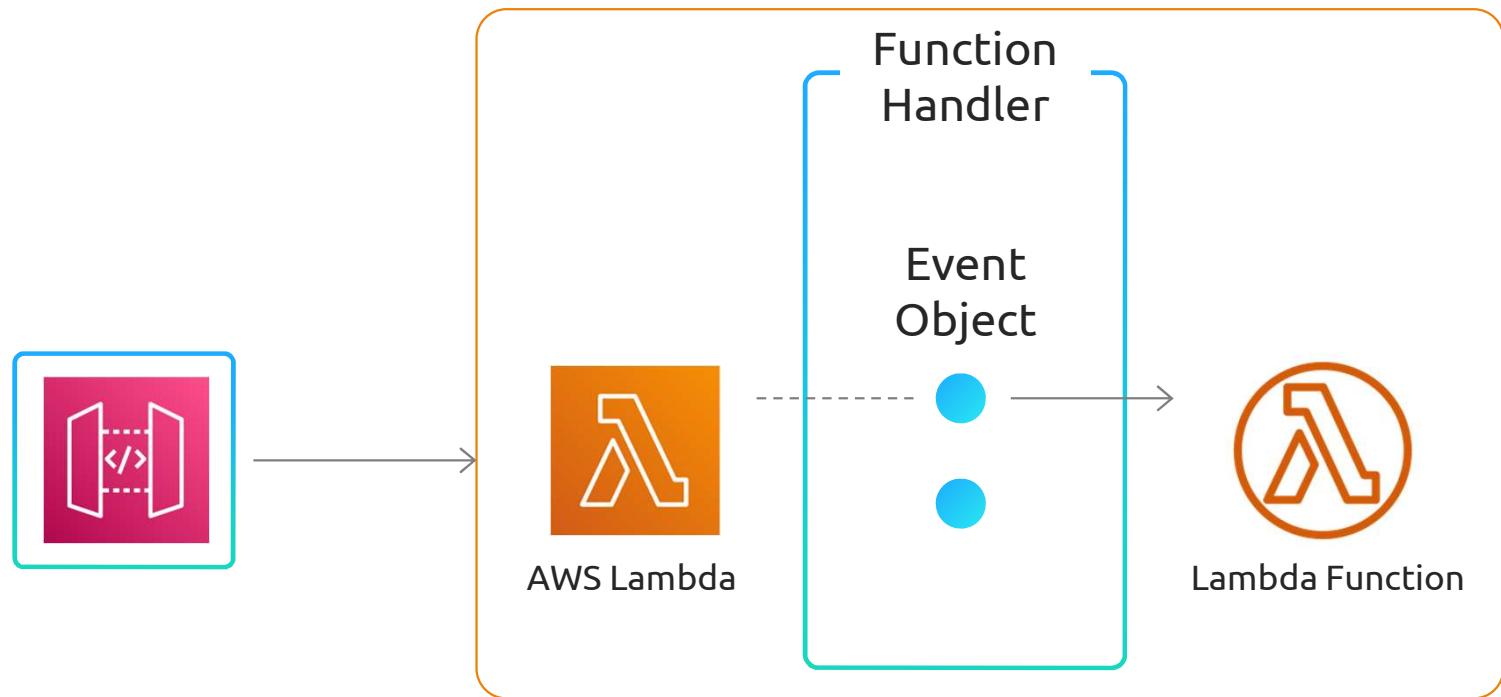


Functions



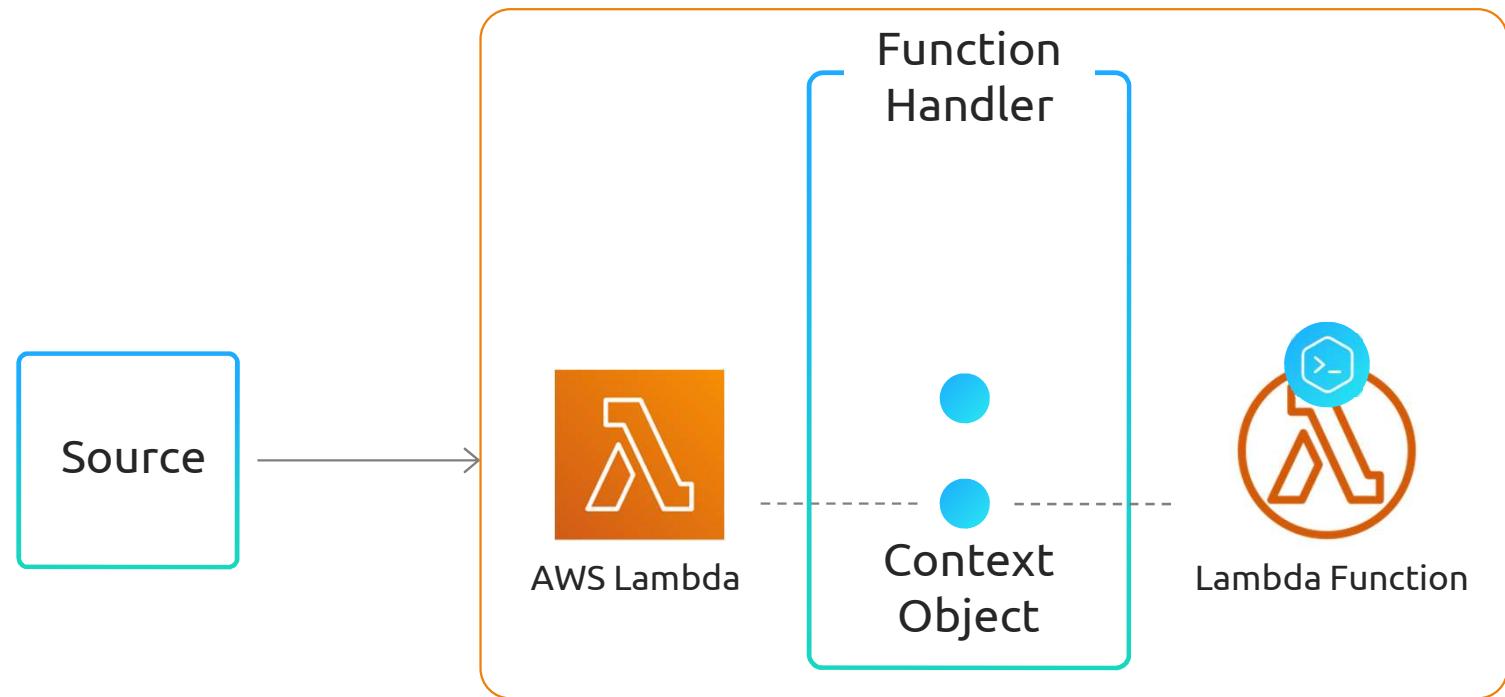


Functions





Functions





```
...                         lambda_function.py

def lambda_handler(event, context):
    message = 'Hello {} {}'.format(event['first_name'], event['last_name'])
    return { 'message' : message }
```

AWS Lambda Developers Guide:

<https://docs.aws.amazon.com/lambda/latest/dg/welcome.html>



Functions

Best Practices:

<https://docs.aws.amazon.com/lambda/latest/dg/best-practices.html>

- Writing Functions
- Testing
- Performance
- Using Lambda with specific services
- Working with streams
- And more...



Summary



AWS Lambda Functions

- Lambda contains a Function that executes the programming code in one of the supported programming languages.
- The user can bring the code without changing its bulk, but a function handler must be included.
- The function handler has two objects: the [event object](#) and the [context object](#).
- [Event Object](#): Allows the event source to pass information to the Lambda function.
- [Context Object](#): Contains information specific to the runtime environment of the programming language used and allows communication within the function.
- The configuration of the code specifics depends on the programming language used.



KodeKloud

Pricing



Pricing

Start Free

Monthly Free Tier:
1 Million Requests
400, 000 Gigabit seconds



Pricing Components

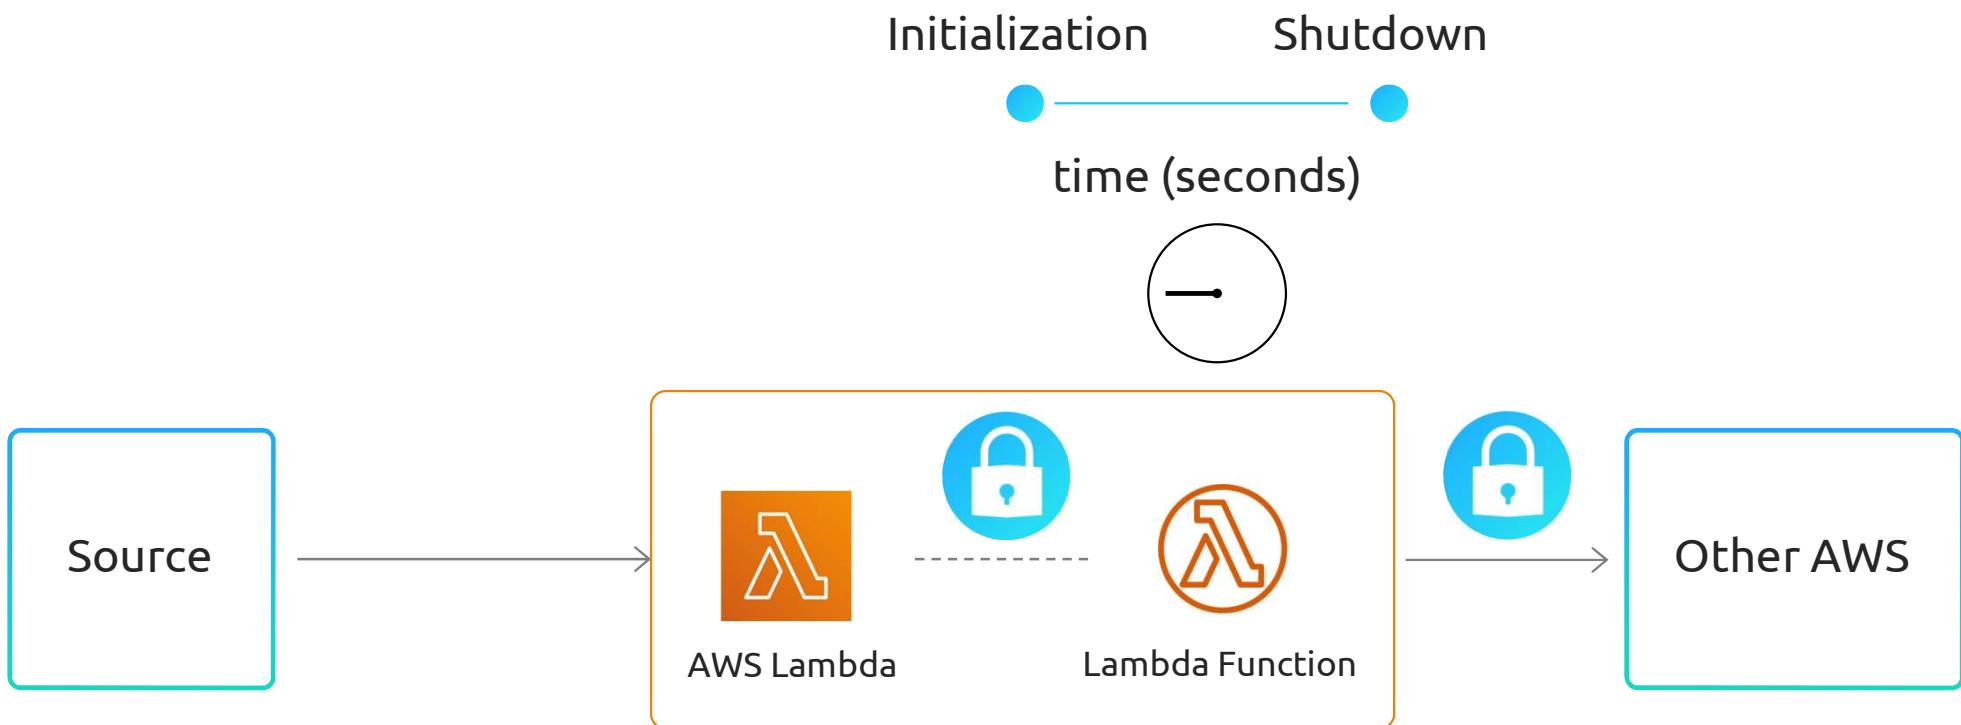
Parts of Lambda Cost

1. Number of Requests
2. Gigabit Seconds (Amount of Time X Amount of Resources)

Lambda Requests

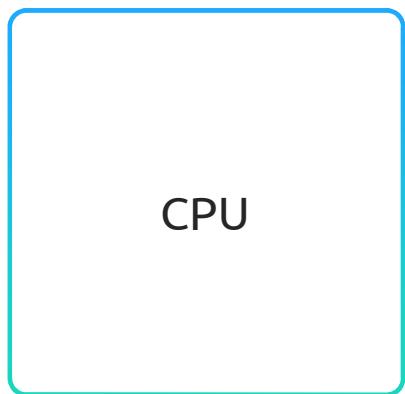


Gigabit Seconds

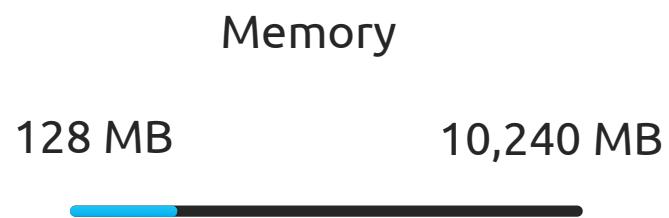




Gigabit Seconds



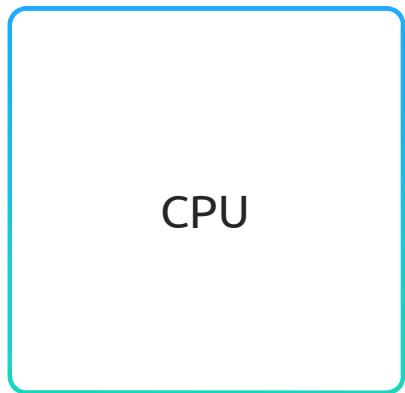
CPU



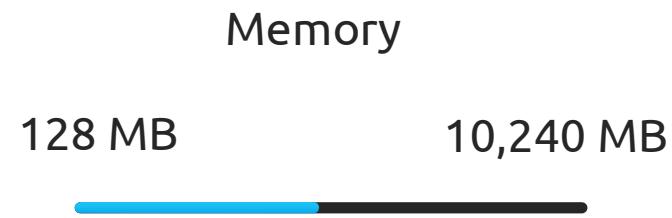
Lambda Function



Gigabit Seconds



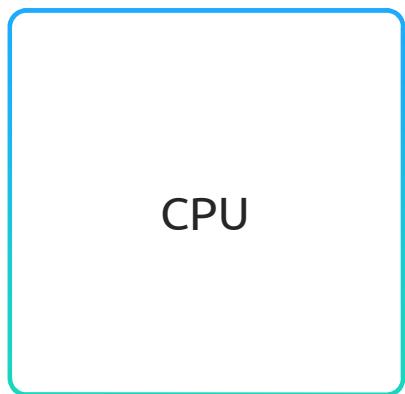
CPU



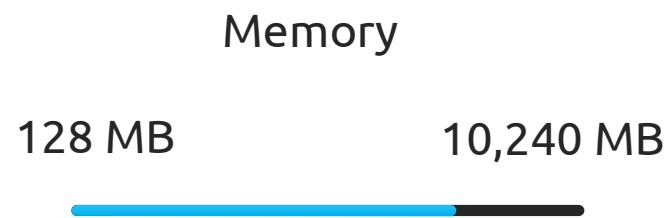
Lambda Function



Gigabit Seconds



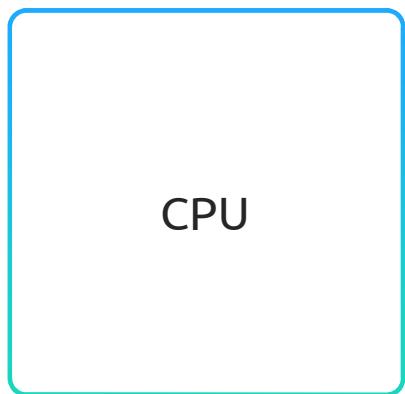
CPU



Lambda Function



Gigabit Seconds



CPU

Memory

128 MB 10,240 MB

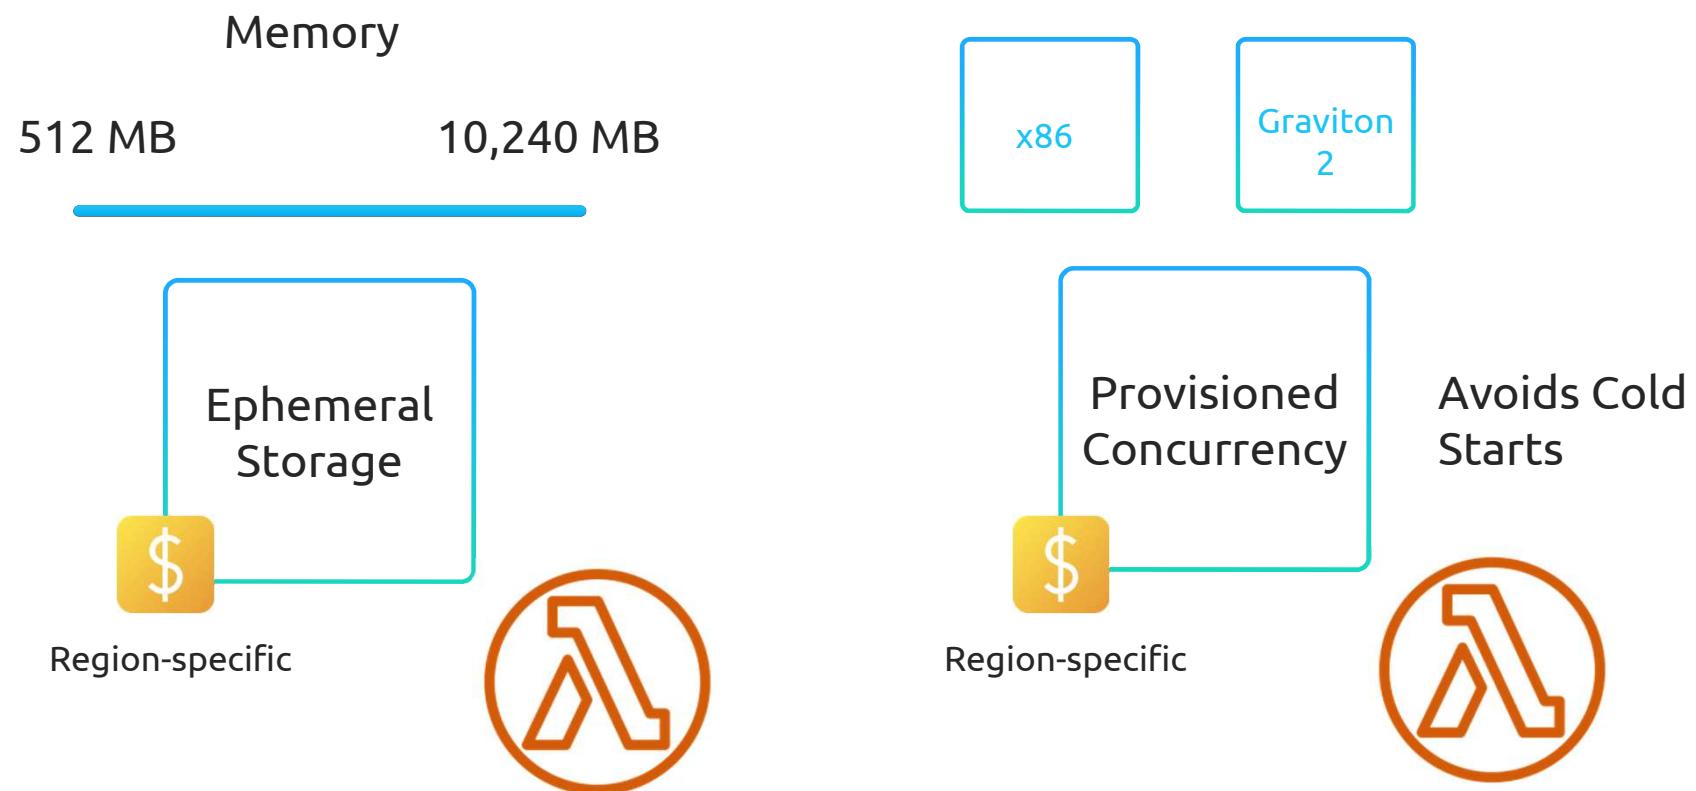


Lambda Function

Pricing Components



Pricing Components





Lambda Pricing Example

AWS Calculator Link:

<https://calculator.aws>



Summary



AWS Lambda Pricing

- Lambda's [free tier](#) allows you to use [1 million requests](#) and [400,000GB seconds per month](#).
- The main components of Lambda pricing are the [number of requests](#) and [Gigabit Seconds](#).
- A Lambda request starts each time it executes the function in response to an event trigger.
- Gigabit seconds are the amount of time your function runs multiplied by the amount of memory it uses.
- The amount of memory you allocate to the function determines the amount of CPU power allocated by Lambda.
- There are [additional costs](#) for including [ephemeral storage](#) and [provisioned concurrency](#) to your Lambda functions.



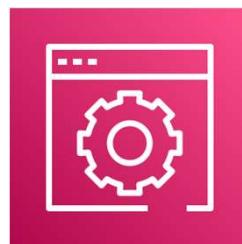
KodeKloud

Prerequisites

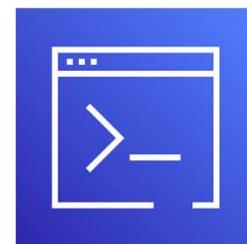


Prerequisites

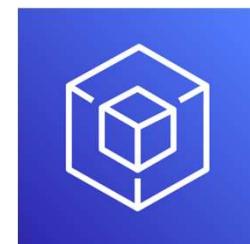
Ways to configure Lambda



AWS
Management Console



AWS Command Line
Interface (CLI)



Software Development
Kits (SDKs)





Prerequisites

Ways to configure Lambda



AWS
Cloudformation

Infrastructure
as Code

JSON or YAML File



AWS
Serverless Application Model



Prerequisites



<https://aws.amazon.com/console/>



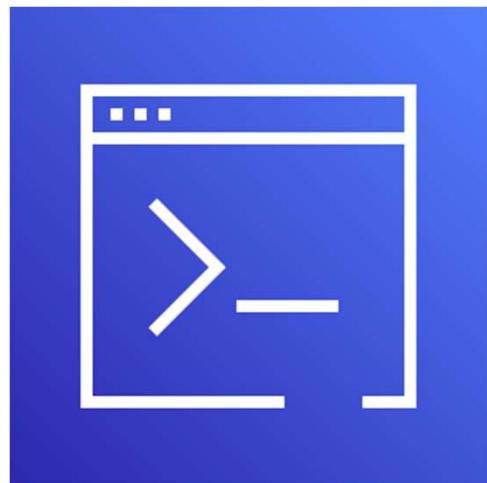
Prerequisites



AWS Command Line
Interface (CLI)



Prerequisites



AWS Command Line Interface (CLI)

Download the application at:

<https://aws.amazon.com/console/>

Supported on:

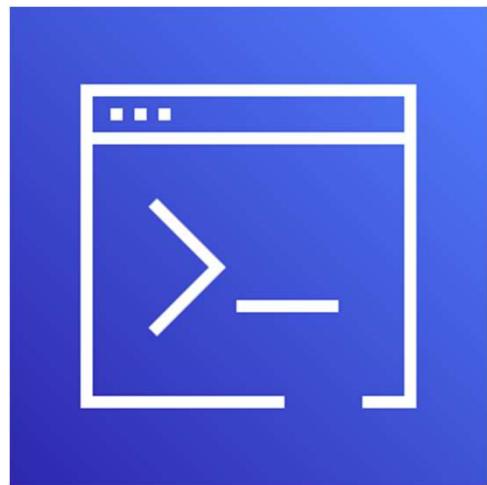
Windows

MacOS

Linux



Prerequisites



AWS Command Line Interface (CLI)

Download the application at:

<https://aws.amazon.com/cli>

Supported on:

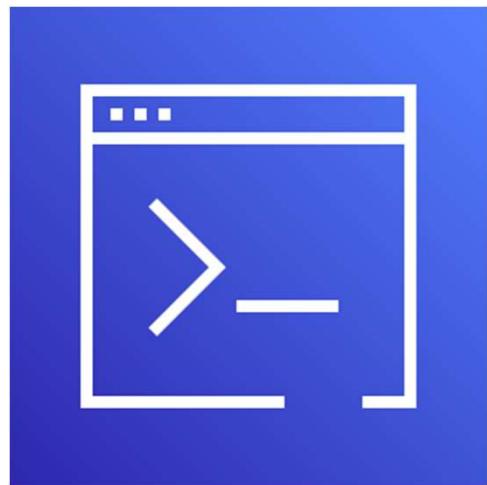
Windows

MacOS

Linux



Prerequisites



AWS Command Line Interface (CLI)

Download the application at:

<https://aws.amazon.com/console/>

Supported on:

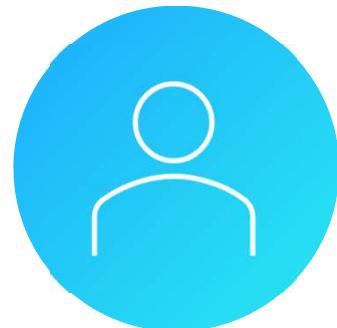
Windows

MacOS

Linux



Prerequisites



Root User Account



Administrative User
Account



Summary

Ways to configure Lambda



AWS
Management Console



AWS Command Line
Interface (CLI)



Software Development
Kits (SDKs)



KodeKloud



Create Basic Function Demo

Environment



AWS CLI



Visual Studio Code
(or equivalent)

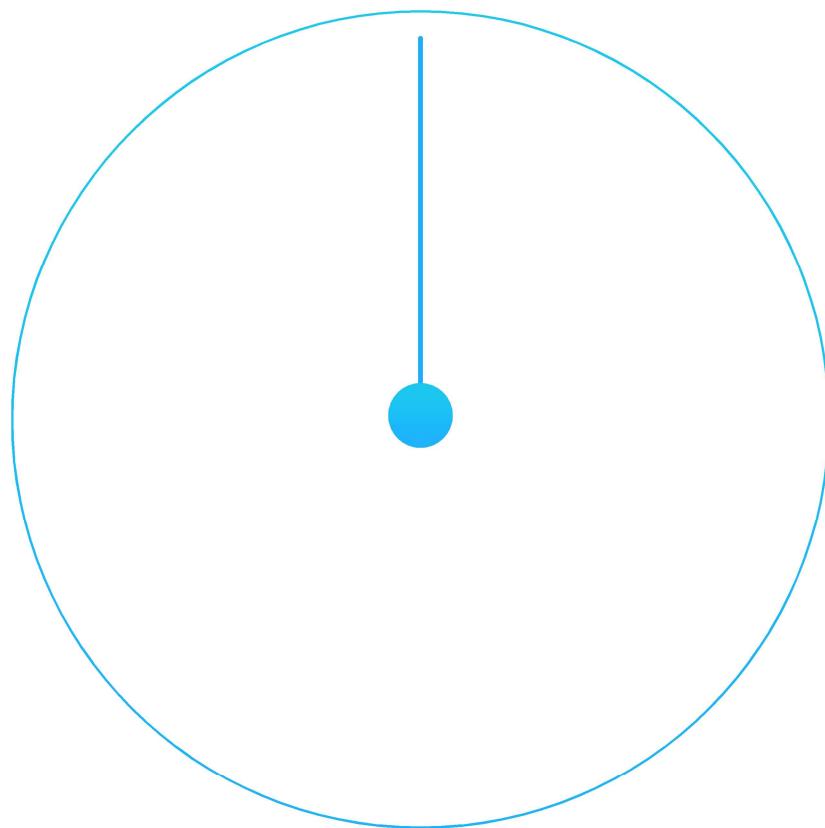


AWS Console

Limitations



Limitations

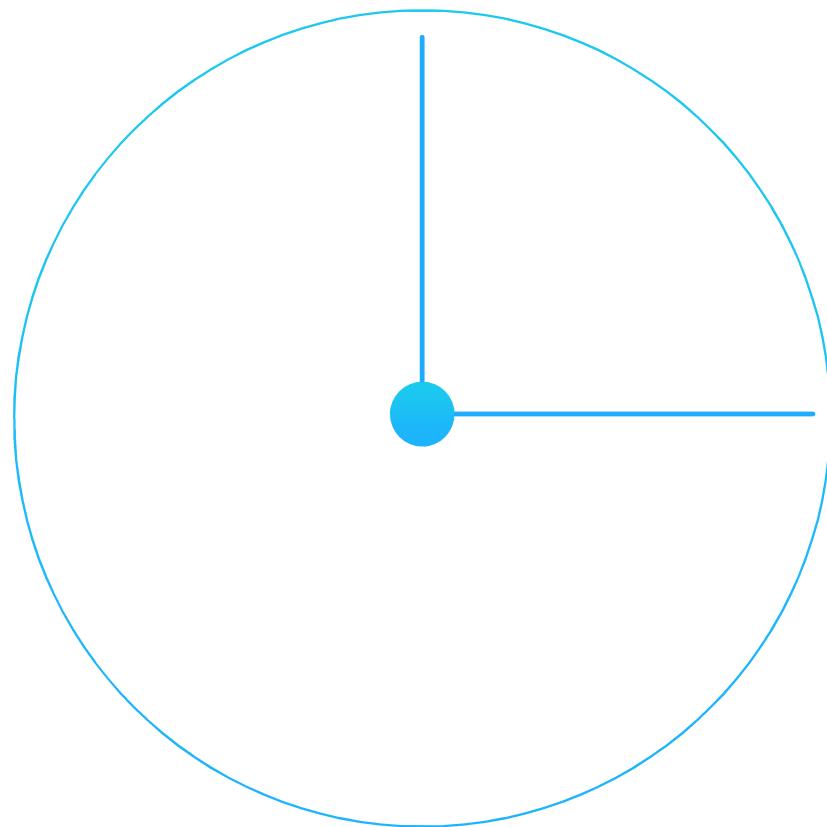




Limitations



Can only run for 15 minutes





Limitations



Can only run for 15 minutes



10 GB RAM Limit



10 GB Storage Limit





Limitations



Can only run for 15 minutes



10 GB RAM Limit



10 GB Storage Limit



1000 Concurrent Execution



1000



Limitations



Can only run for 15 minutes



10 GB RAM Limit



10 GB Storage Limit



1000 Concurrent Execution



3000

Region-specific

Limitations



Can only run for 15 minutes



10 GB RAM Limit



10 GB Storage Limit



1000 Concurrent Execution

For more information on Limits and Restrictions visit:

<https://docs.aws.amazon.com/lambda/latest/dg/gettingstarted-limits.html>



Summary



Limitations

- Lambda function can only run for **15 minutes maximum**.
- Functions must adhere to a **10 GB RAM** and **10 GB storage limit**.
- **Ephemeral storage** is the **temporary space** for functions to read and write data.
- Lambda **automatically scales** based on load and can handle up to **1000 concurrent executions** running at the same time.
- There is a **burst feature** that allows temporary increase in concurrent executions up to **3000 in some regions**.
- Requesting an increase to the default 1000 concurrent execution limit is possible but only granted by AWS if required and based on your use case.

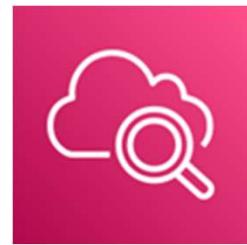


KodeKloud

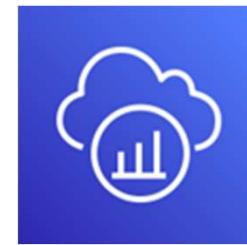
Monitoring



Monitoring

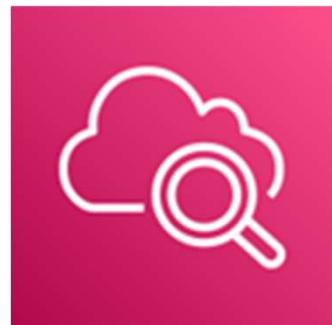


Amazon
Cloudwatch



X-Ray
Services

Amazon Cloudwatch



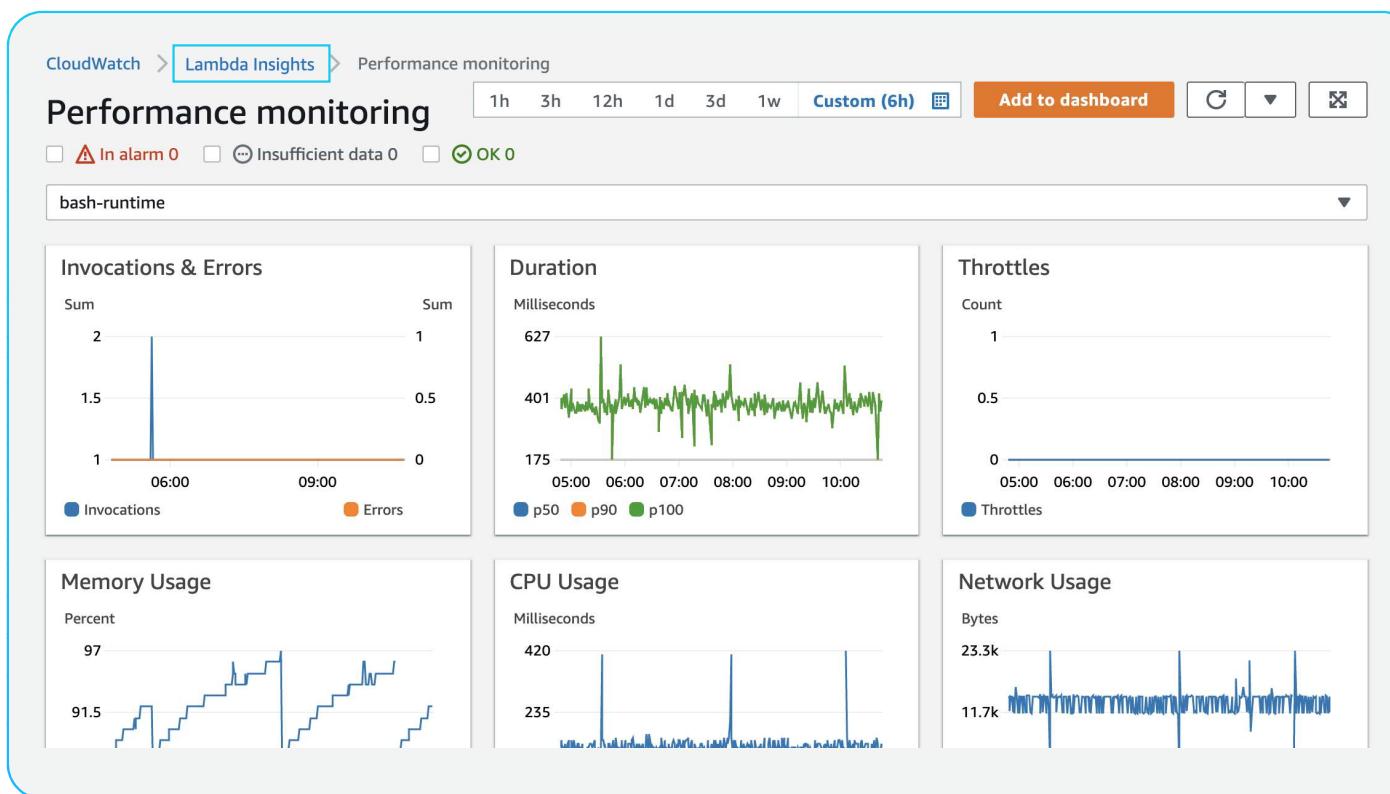
Automatically monitors:

- Number of requests
- Duration per request
- Number of requests resulting to error

Additional Metrics Available:

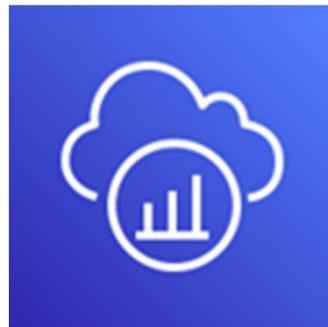
- Throttles due to concurrency limit
- Delay pulling data from stream sources referred to as “Iterator Age”
- Number of events in dead letter queue
- Detailed concurrent execution metrics

Lambda Insights





AWS X-ray



Monitoring and troubleshooting tool

- Visual mapping
- Identify performance bottlenecks and errors
- Trace Lambda function path



Monitoring

Transmission Control Protocol
(TCP)

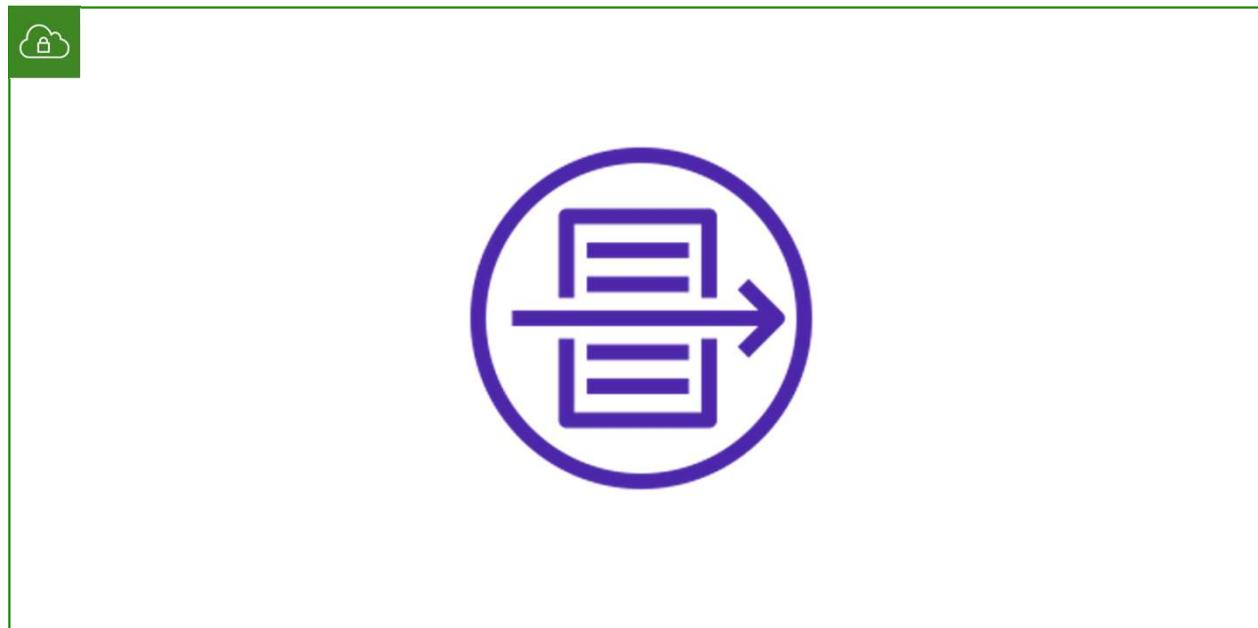
Internet Protocol
(IP)



Flow Logs



Monitoring





Summary



Monitoring

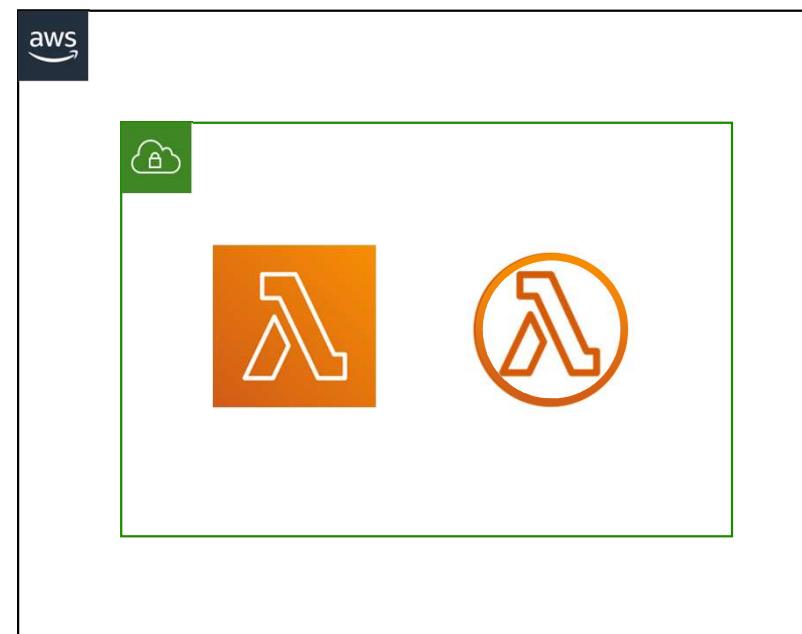
- Lambda integrates with CloudWatch and X-Ray for monitoring and troubleshooting.
- CloudWatch monitors the [number of invocations](#), duration of each request, errors, and other metrics.
- Lambda Insights is an additional [monitoring extension](#) that [provides aggregate information](#) for all functions in your account or region.
- X-Ray is a [monitoring and troubleshooting tool](#) that allows you to [view requests](#) as they travel through your application and [identify bottlenecks and errors](#).
- TCP/IP and network traffic information are unavailable by default for Lambda but may be available through Flow Logs if Lambda runs inside a Virtual Private Cloud (VPC).



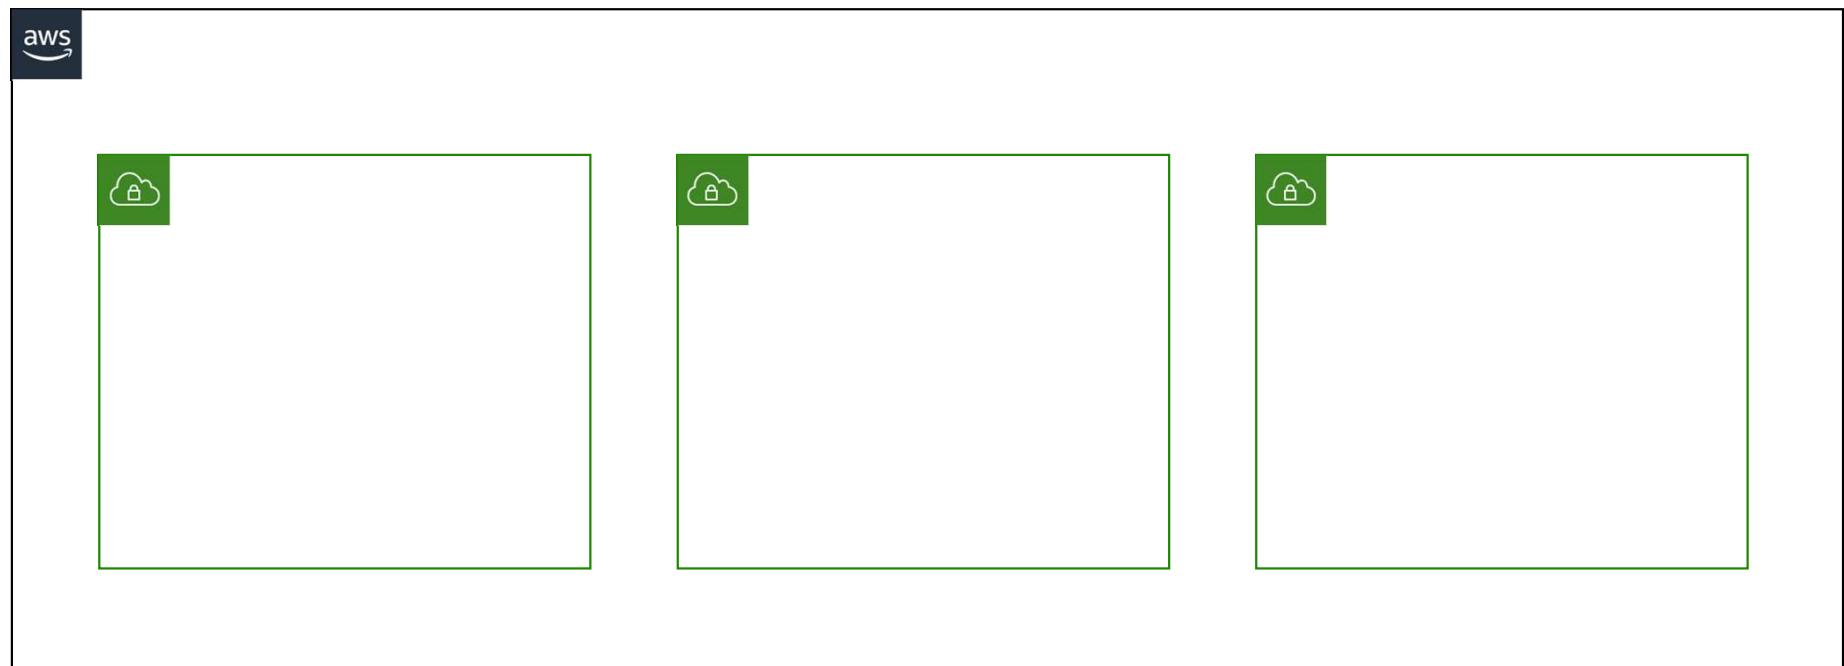
KodeKloud

Lambda Networking

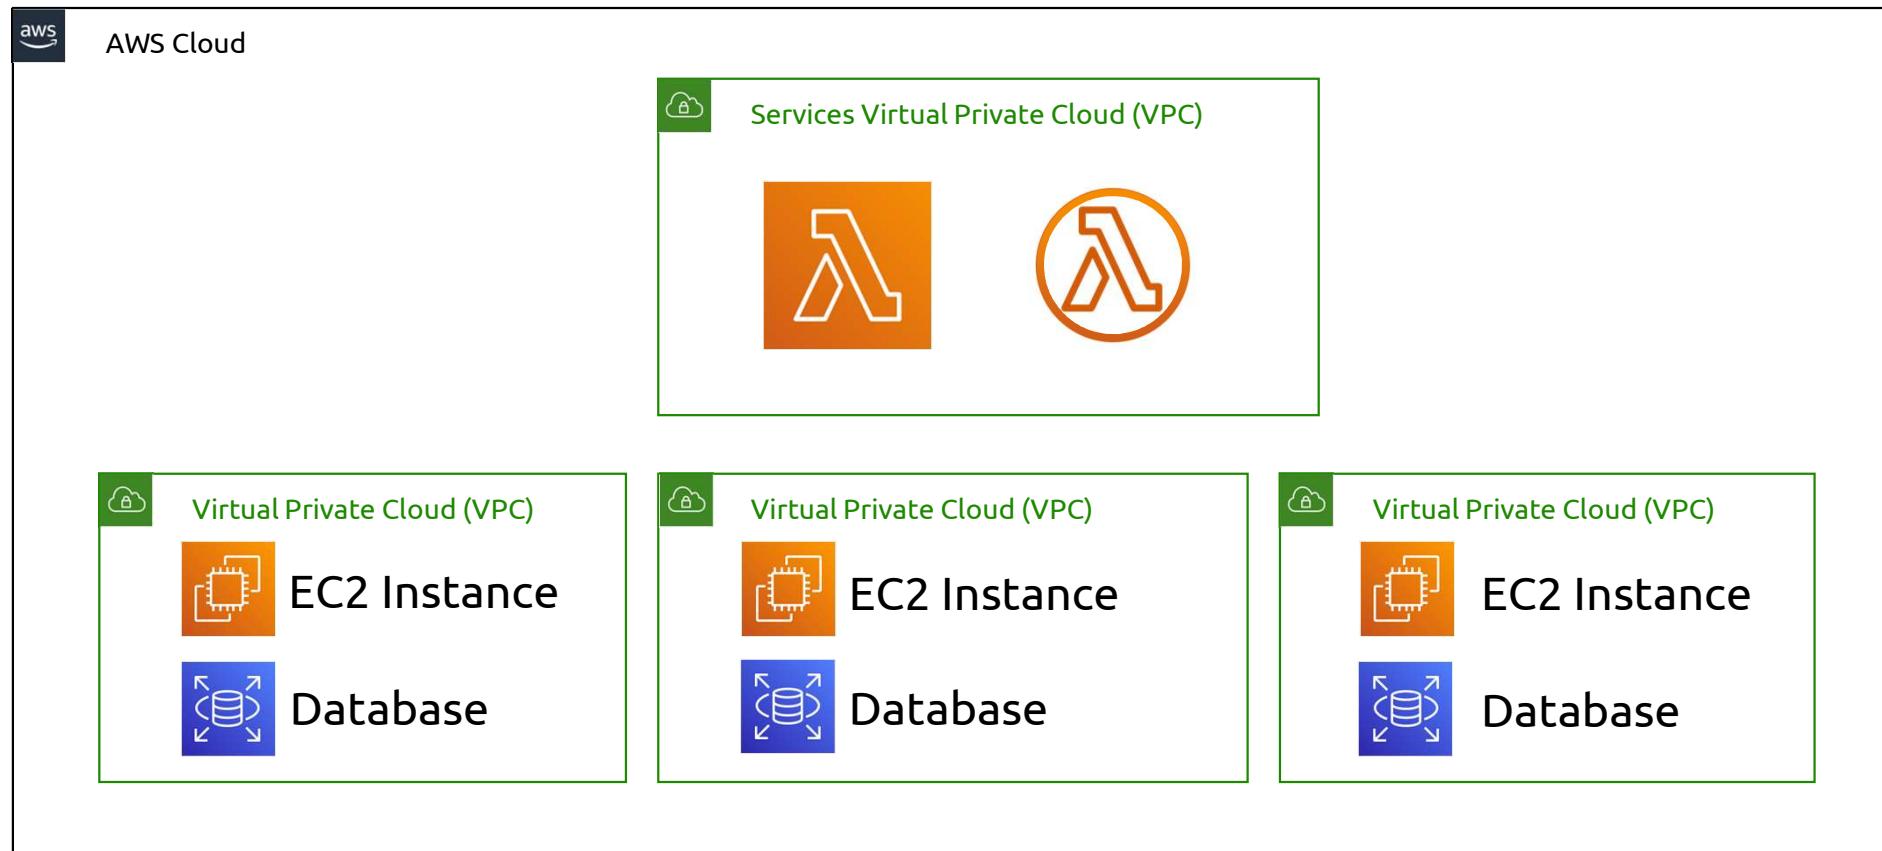
Lambda Networking



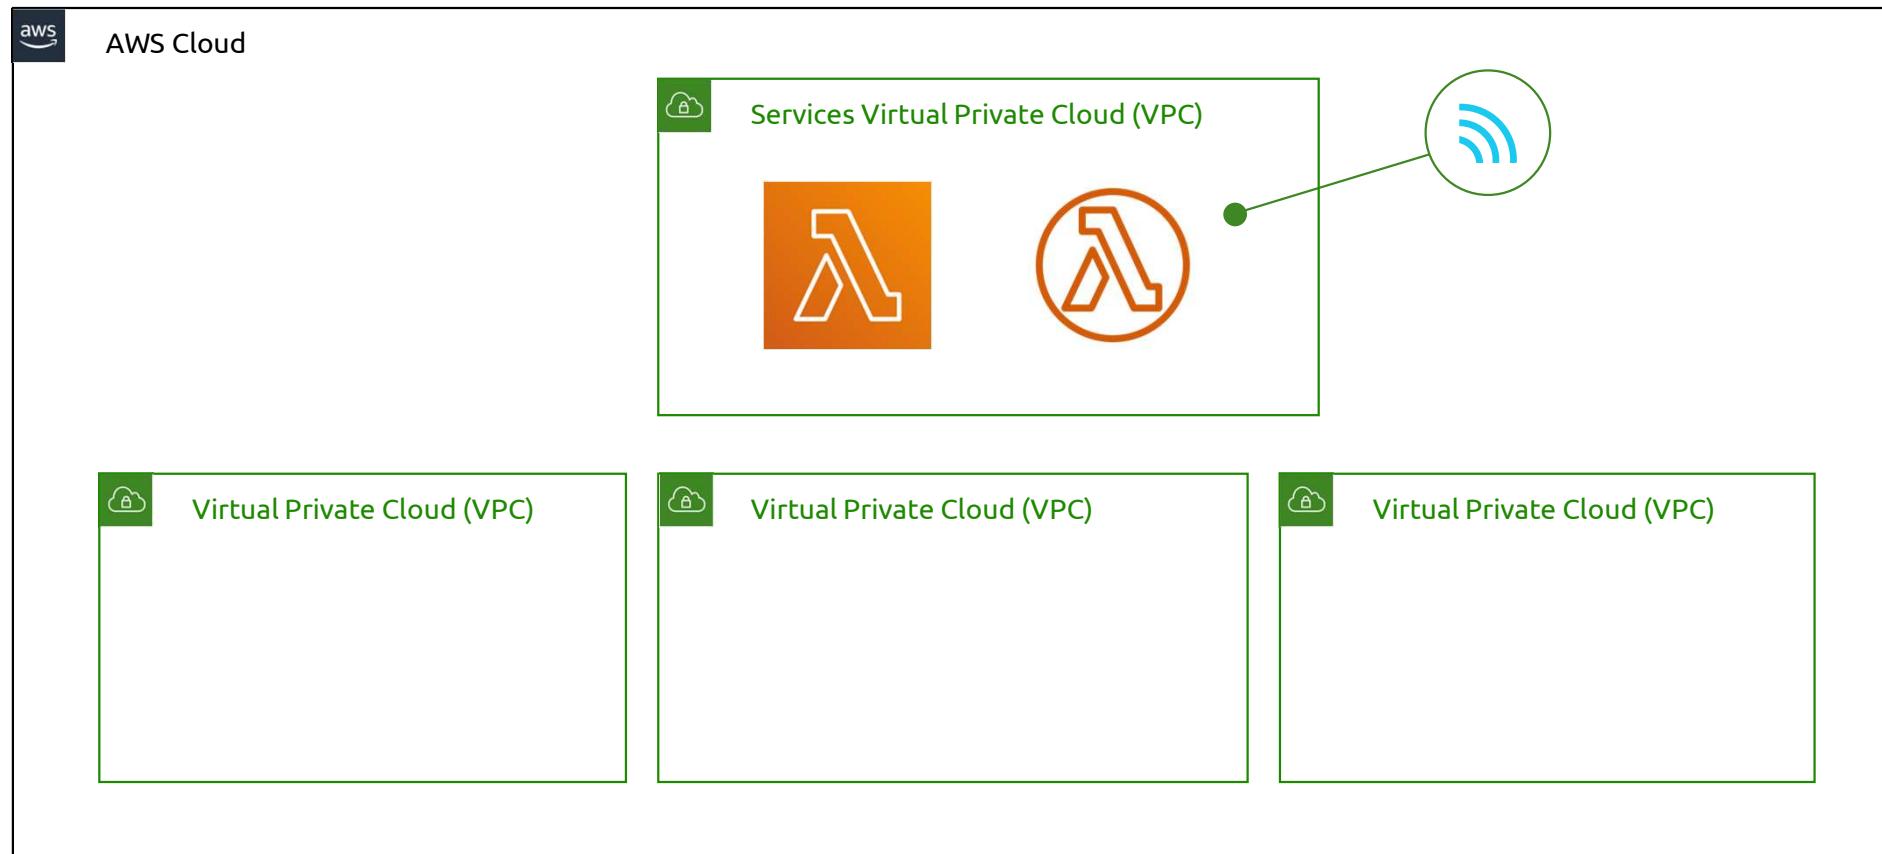
Lambda Networking



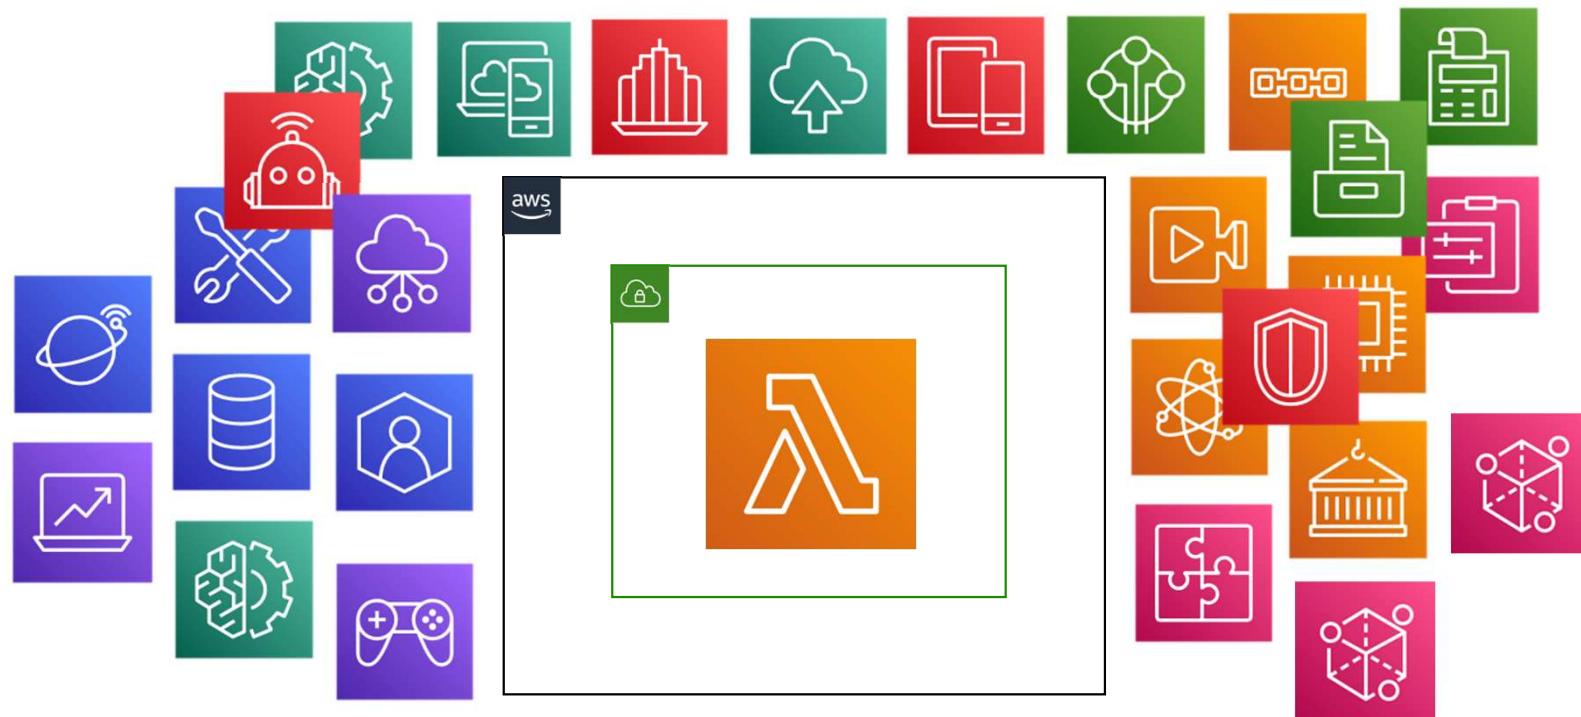
Lambda Networking



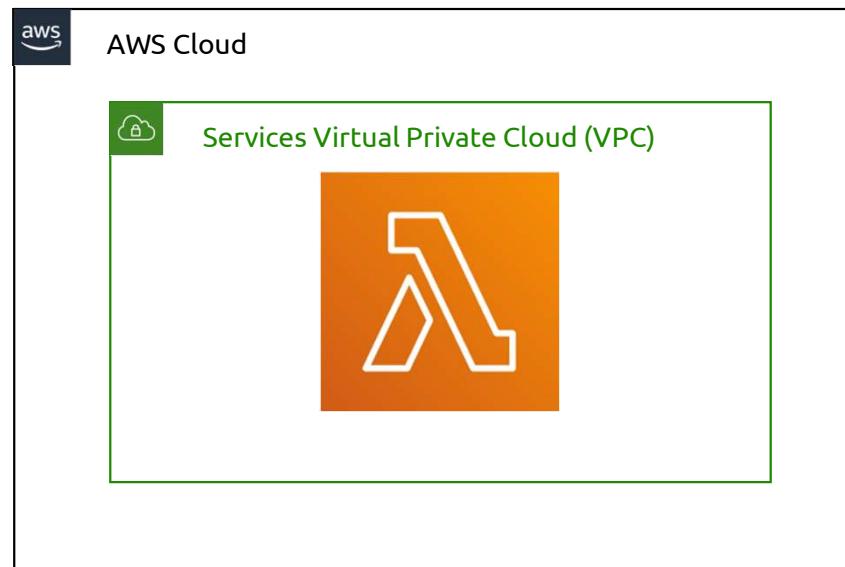
Lambda Networking



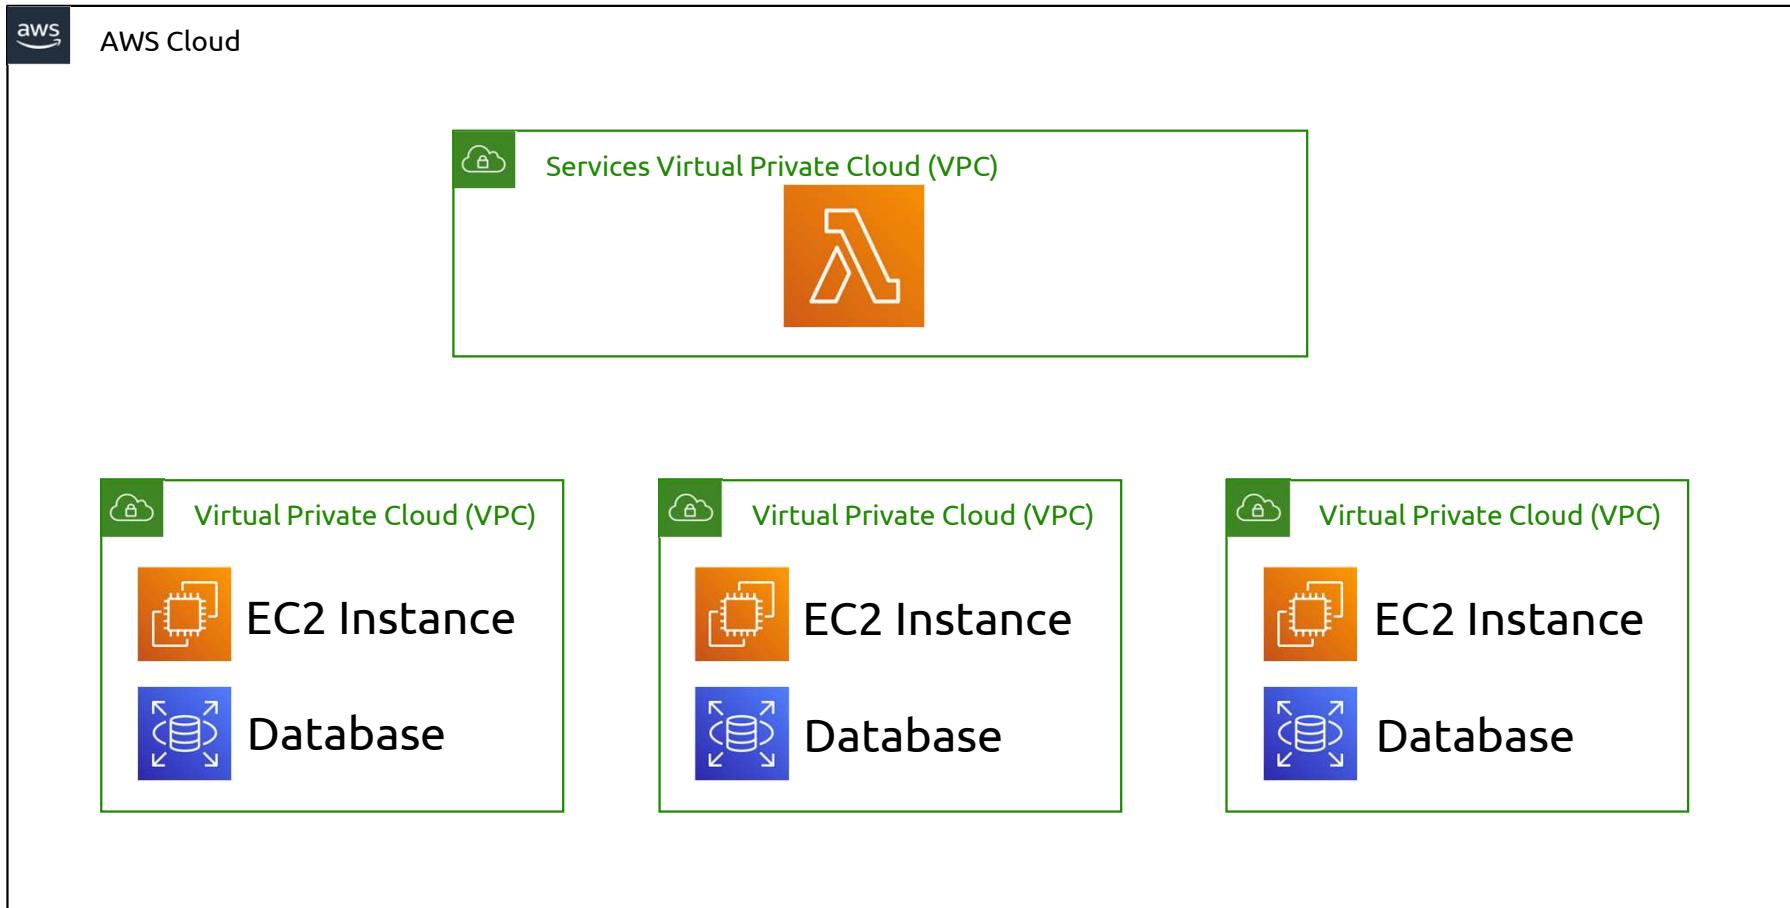
Lambda Networking



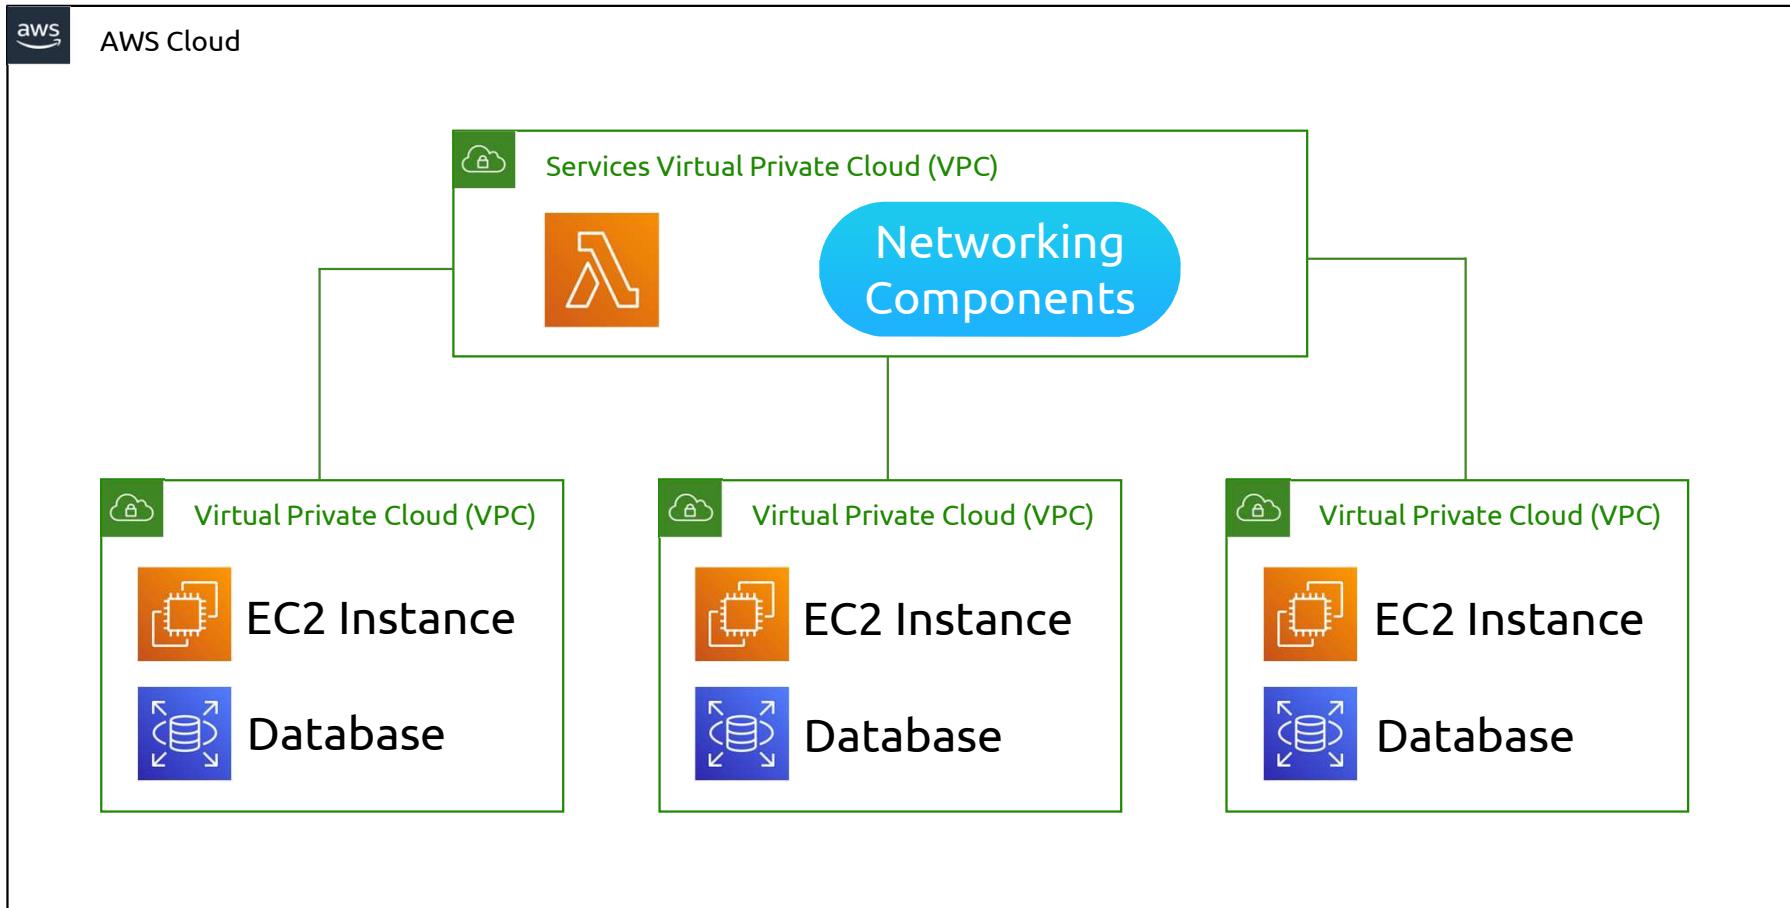
► Lambda Networking



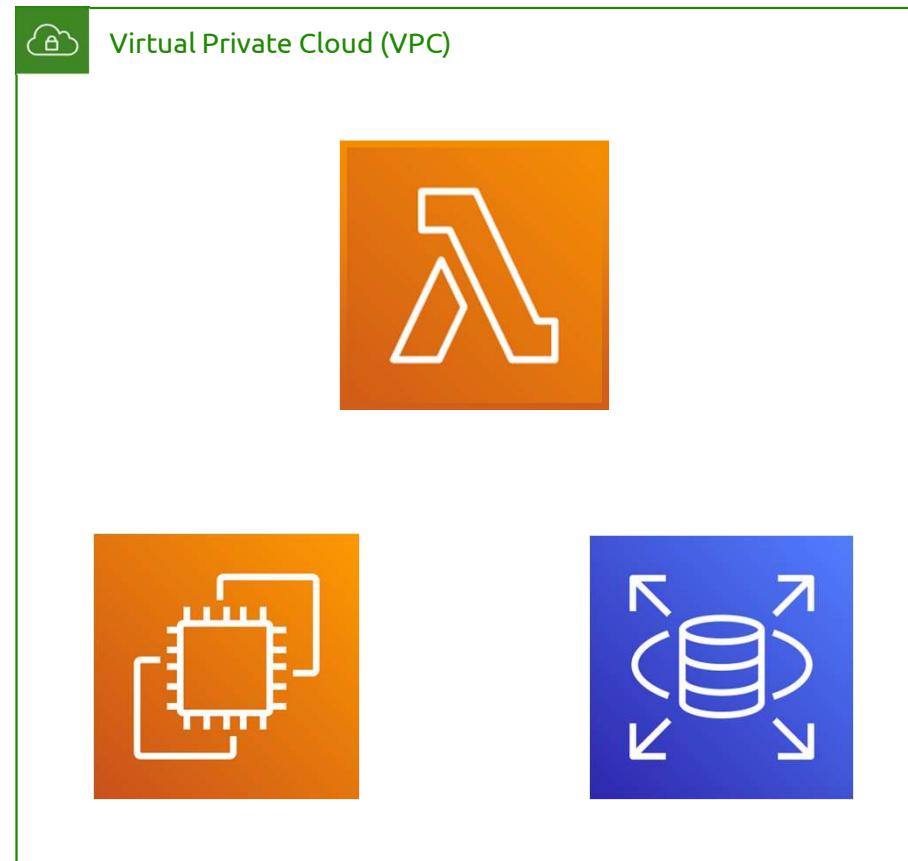
Lambda Networking



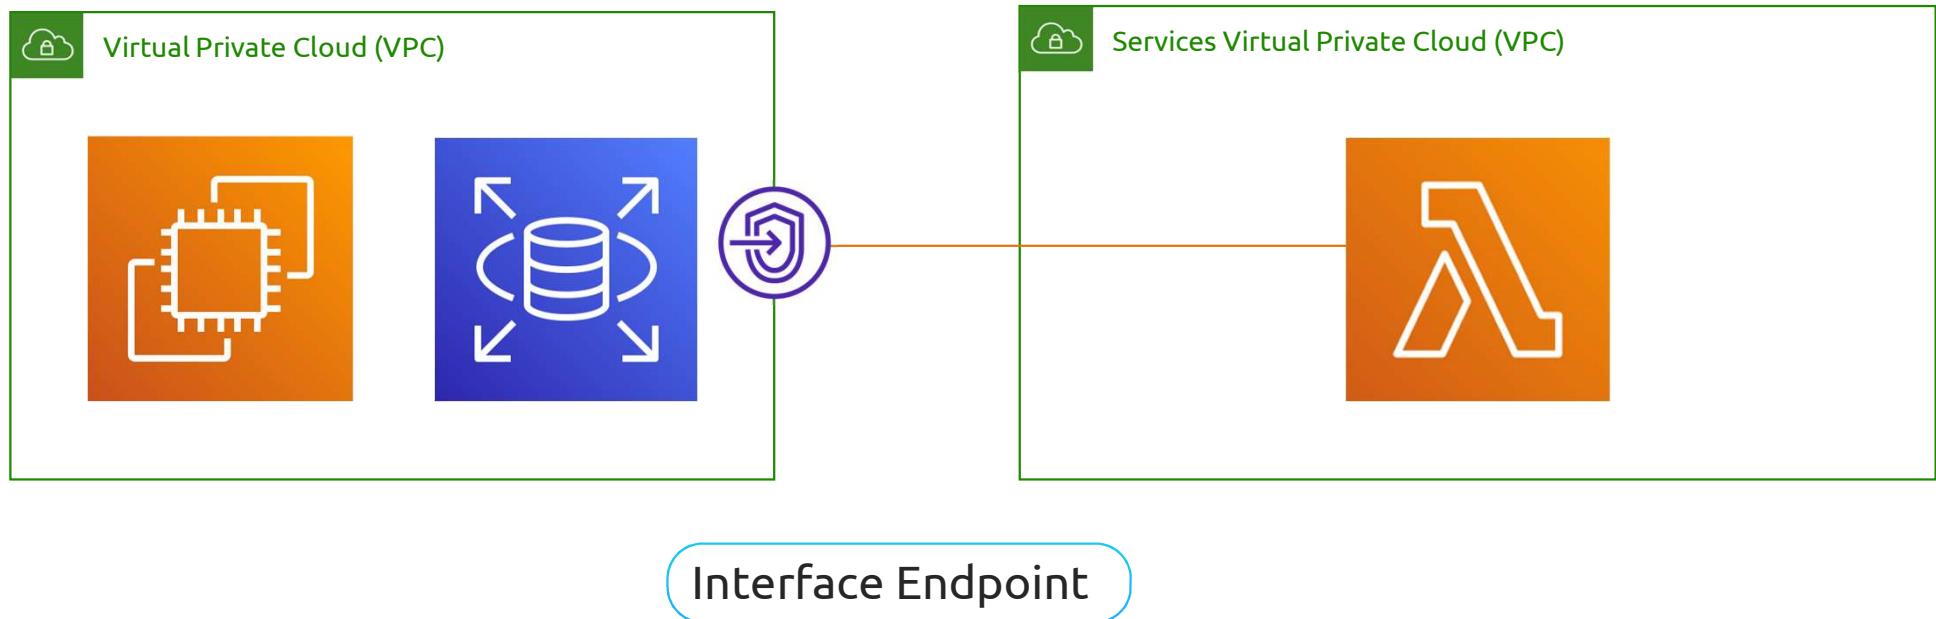
Lambda Networking



Lambda Networking



Lambda Networking



Interface Endpoint

Lambda Networking

▼ Advanced settings

Enable Code signing Info
Use code signing configurations to ensure that the code has been signed by an approved source and has not been altered since signing.

Enable function URL Info
Use function URLs to assign HTTP(S) endpoints to your Lambda function.

Enable tags Info
A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources, track your AWS costs, and enforce attribute-based access control.

Enable VPC Info
Connect your function to a VPC to access private resources during invocation.

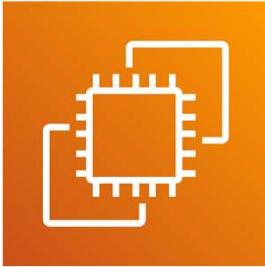
VPC
Choose a VPC for your function to access.
 ▾ C

Cancel Create function

Lambda Networking



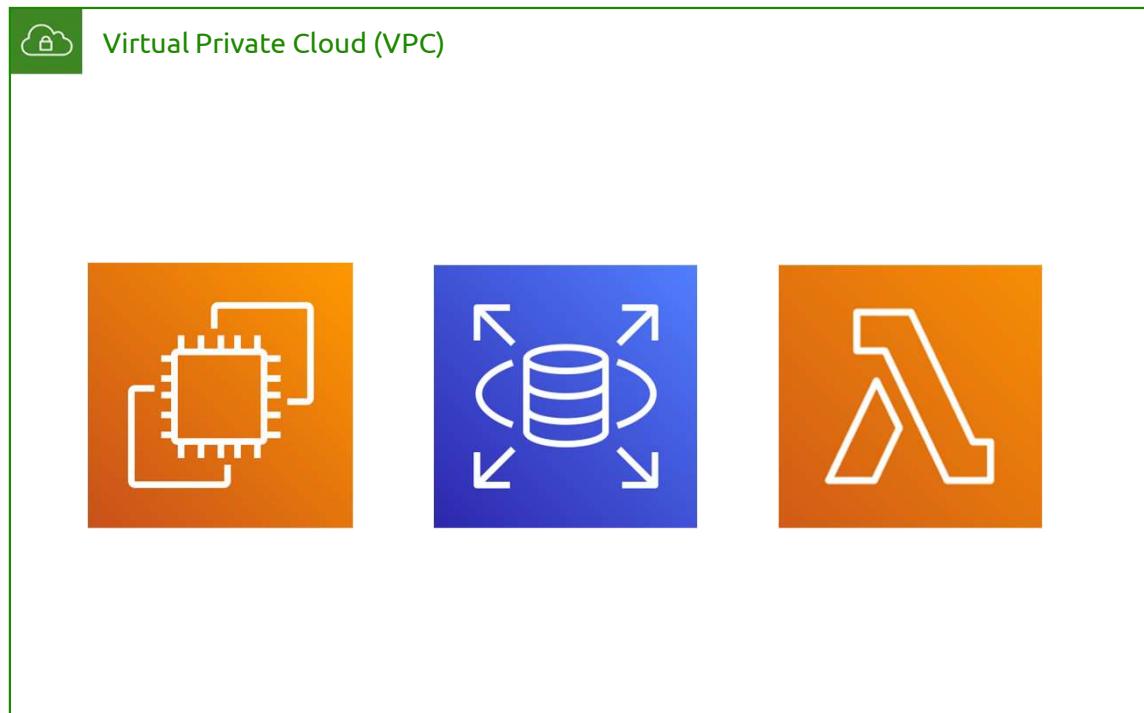
Virtual Private Cloud (VPC)



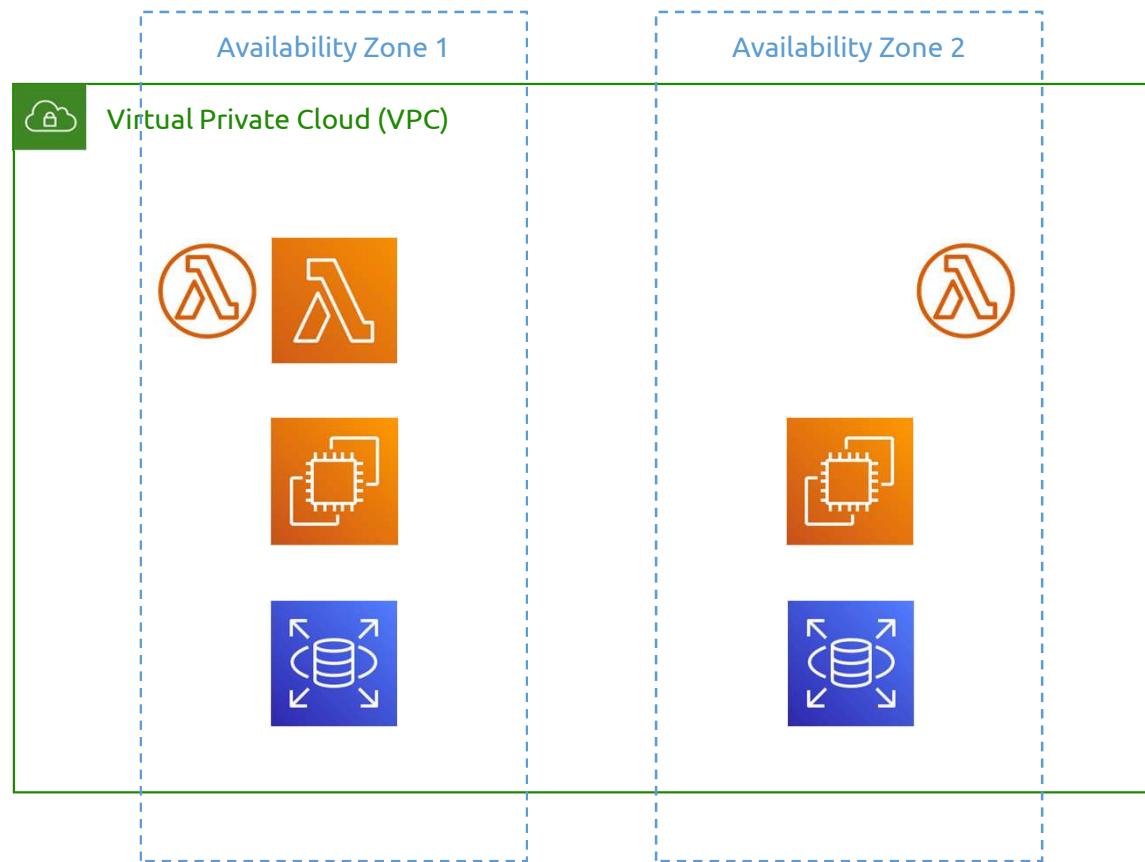
Services Virtual Private Cloud (VPC)



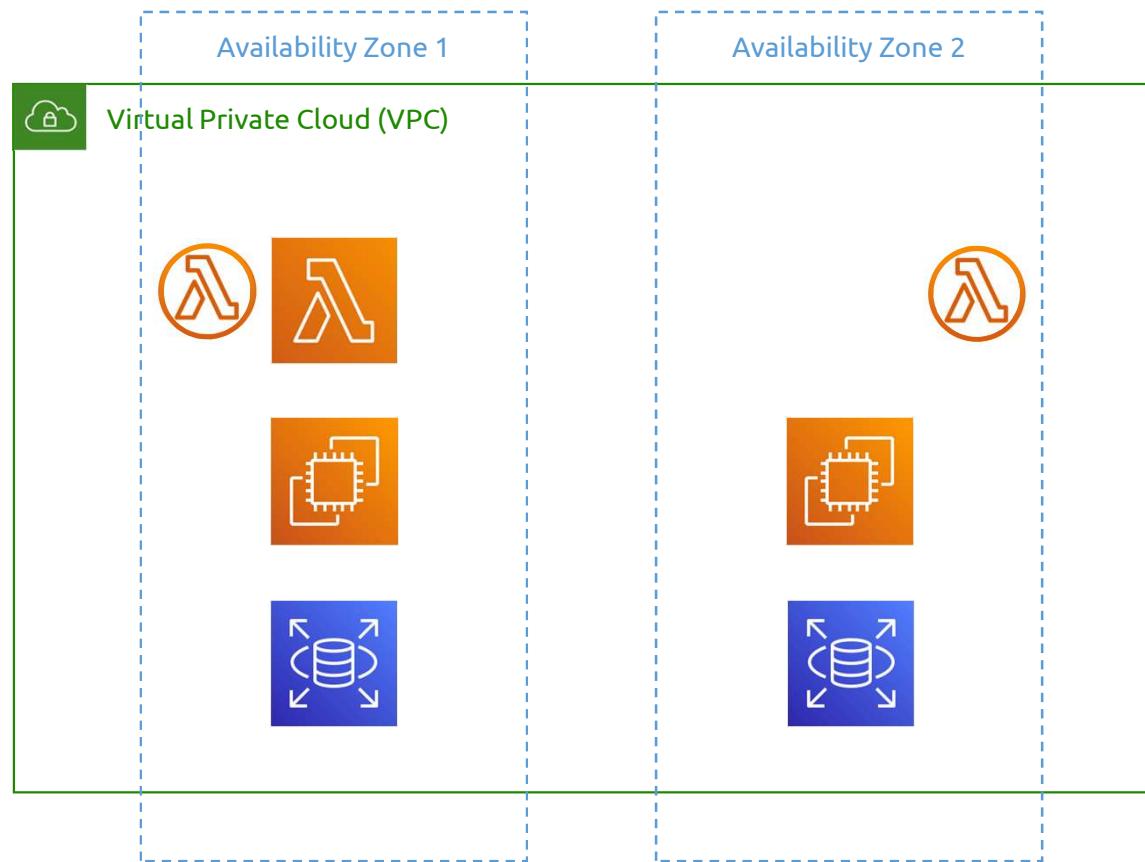
Lambda Networking



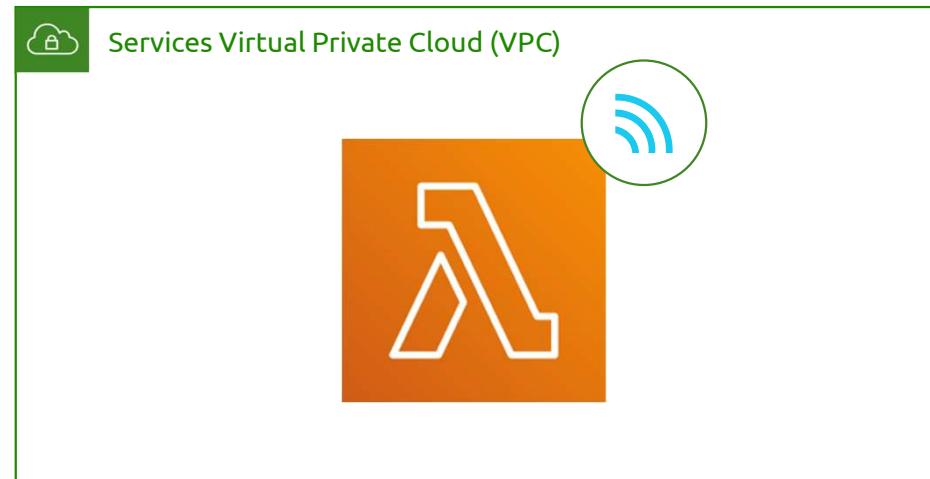
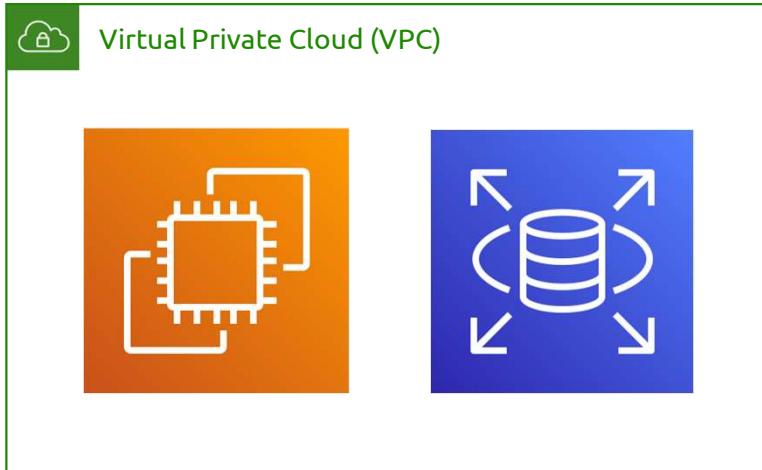
Lambda Networking



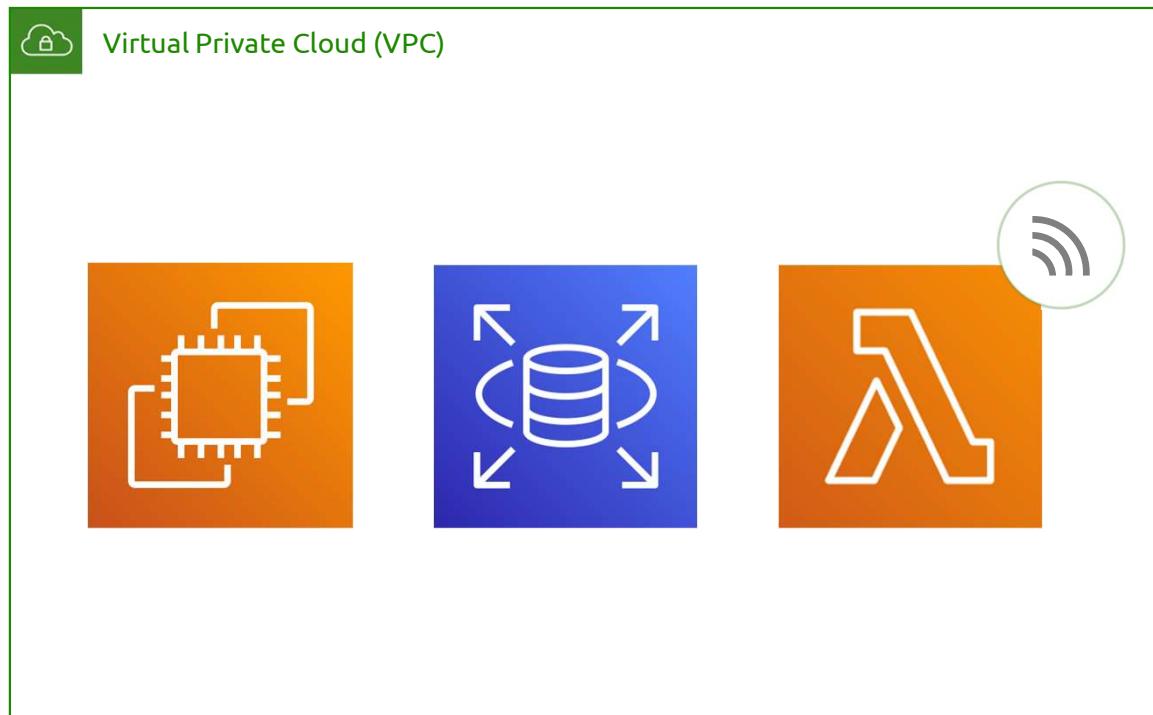
Lambda Networking



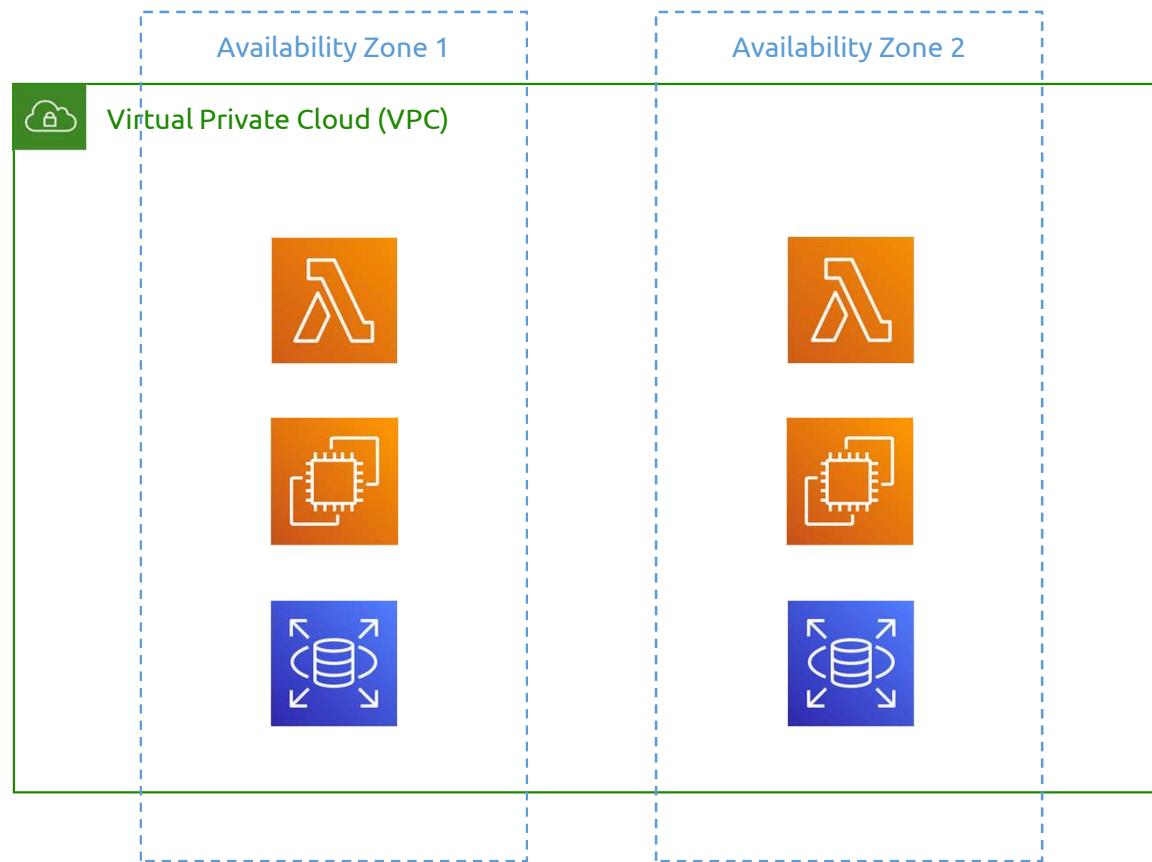
Lambda Networking



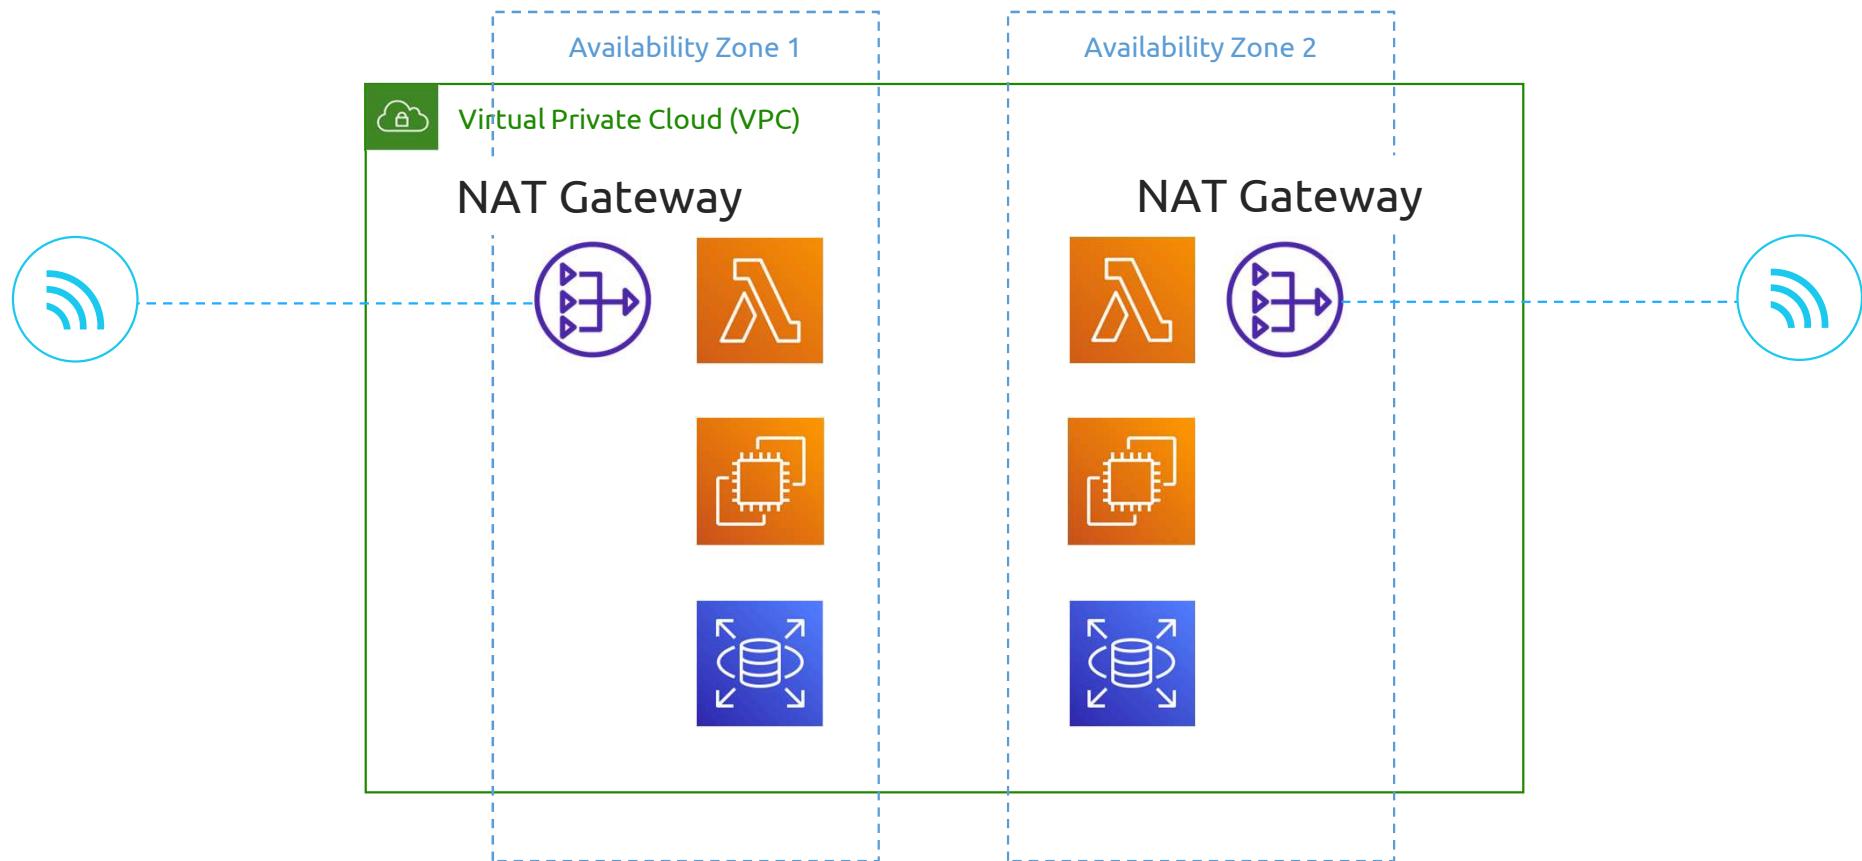
Lambda Networking



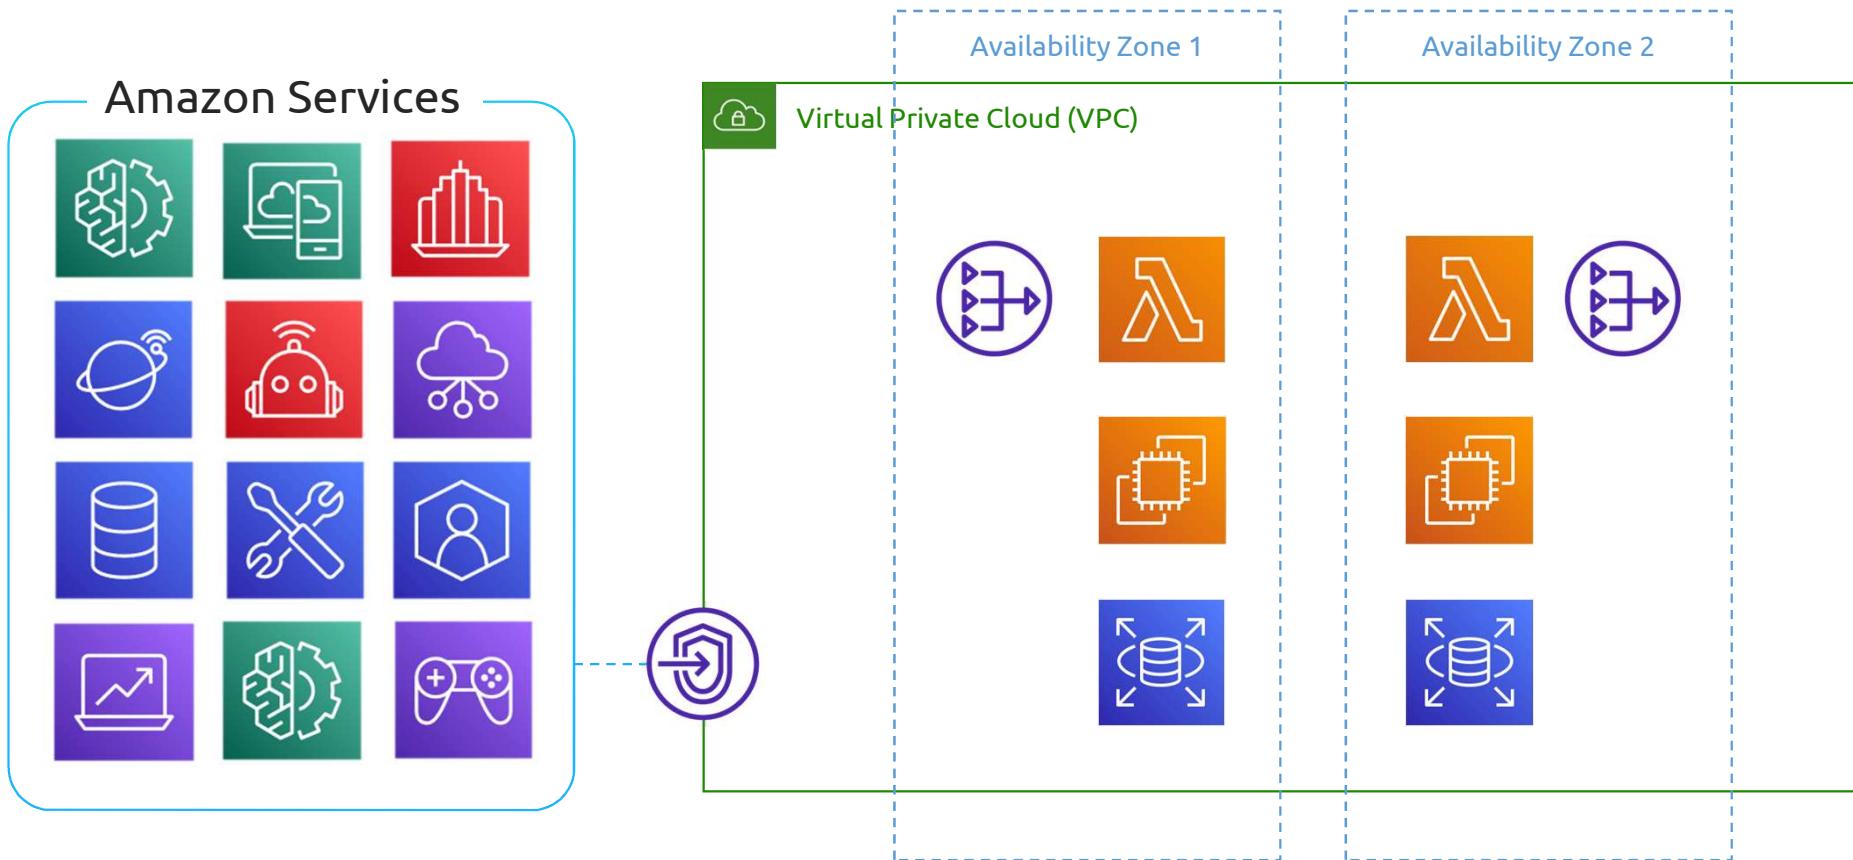
Lambda Networking



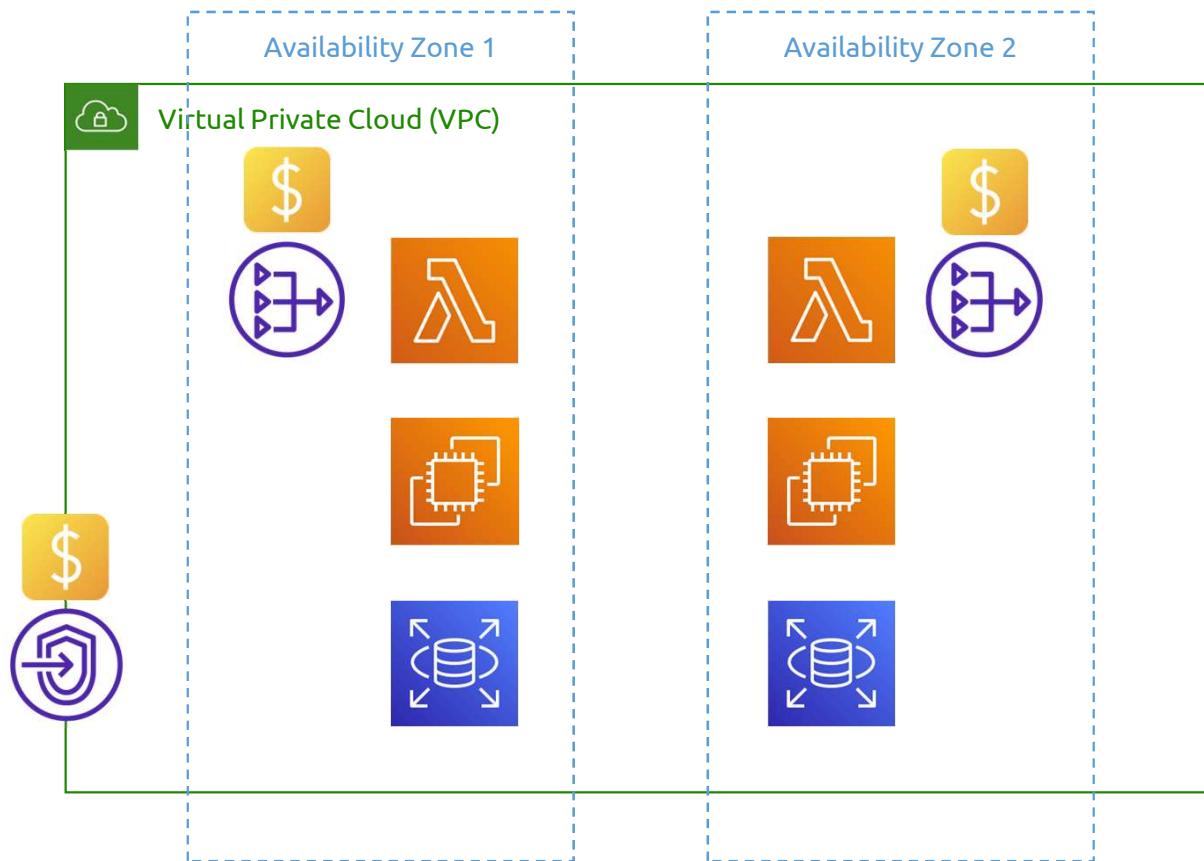
Lambda Networking



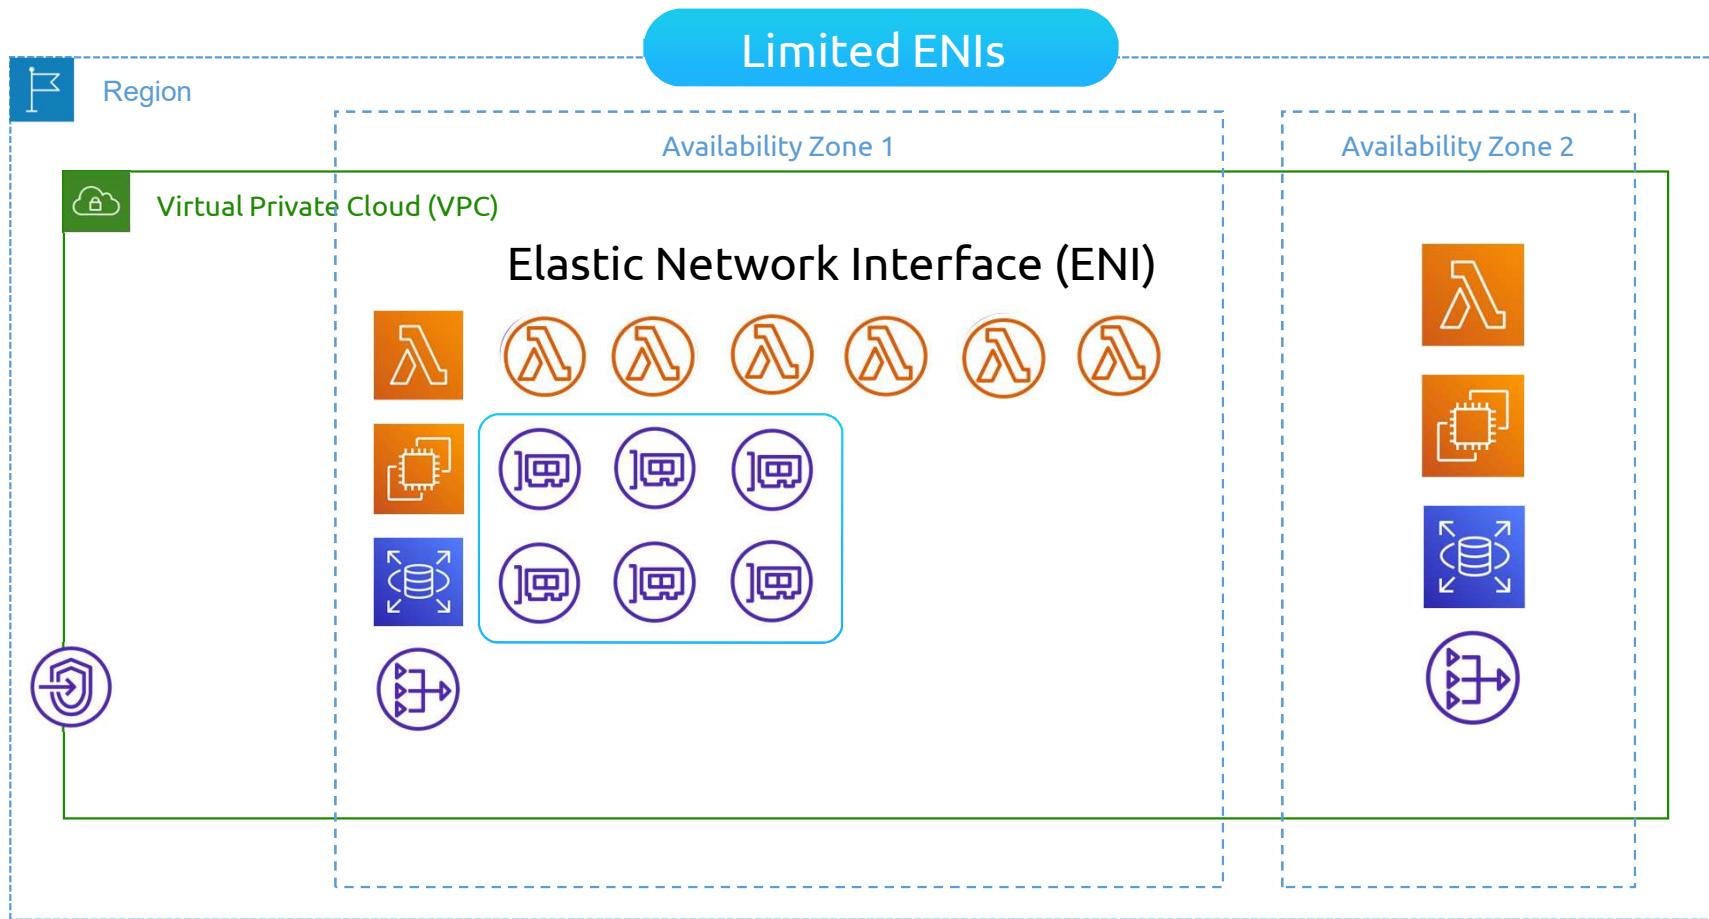
Lambda Networking



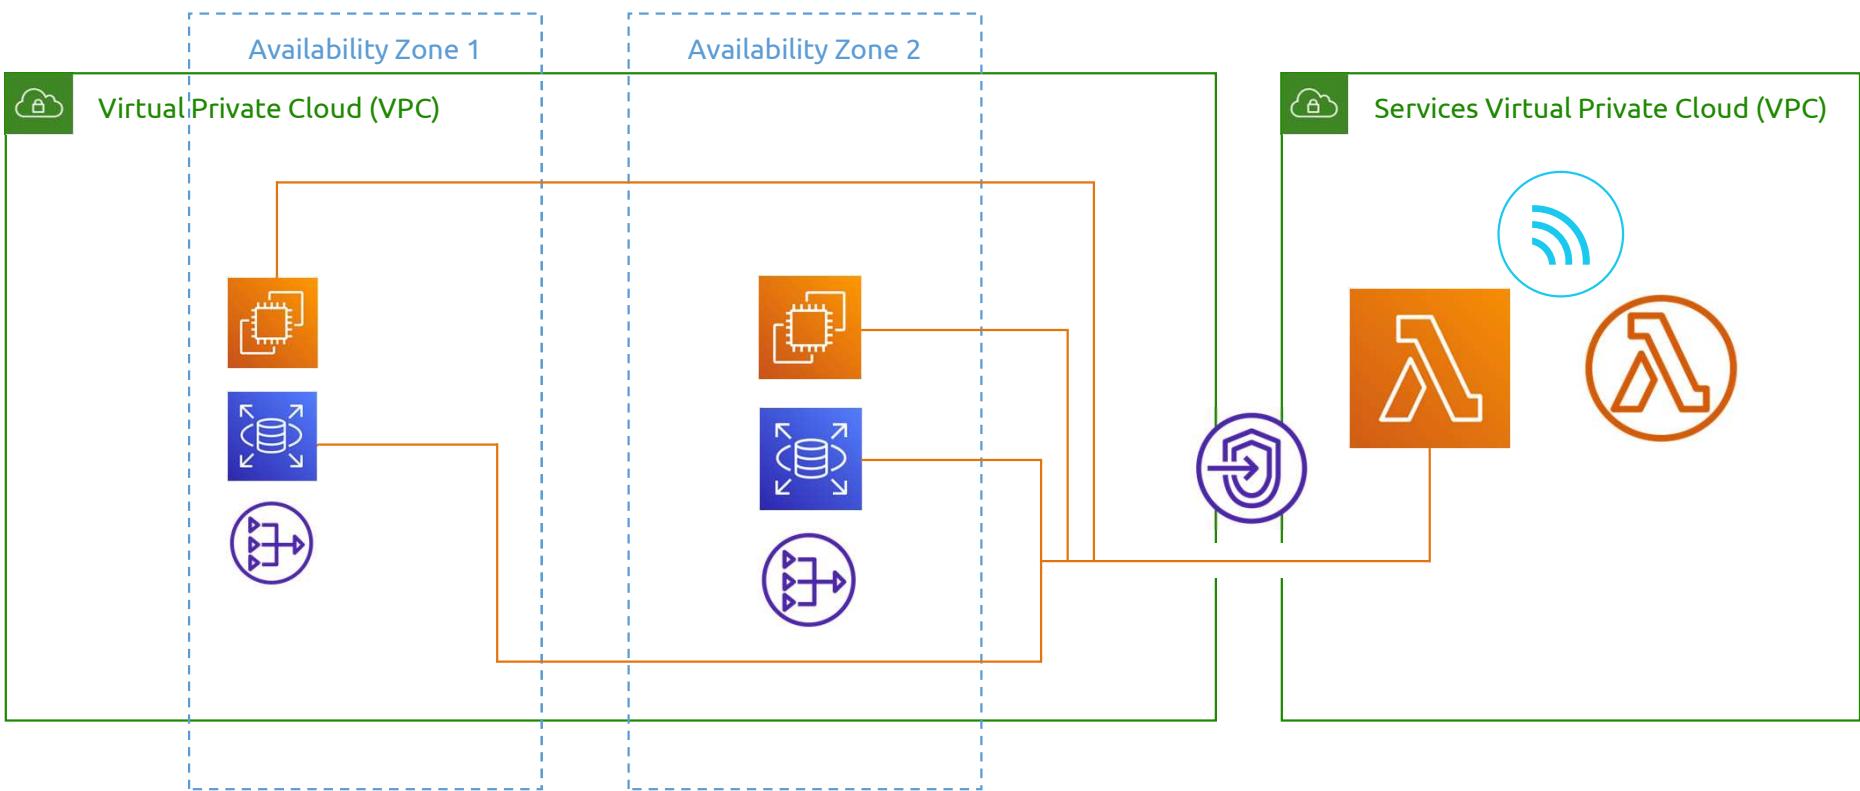
Lambda Networking



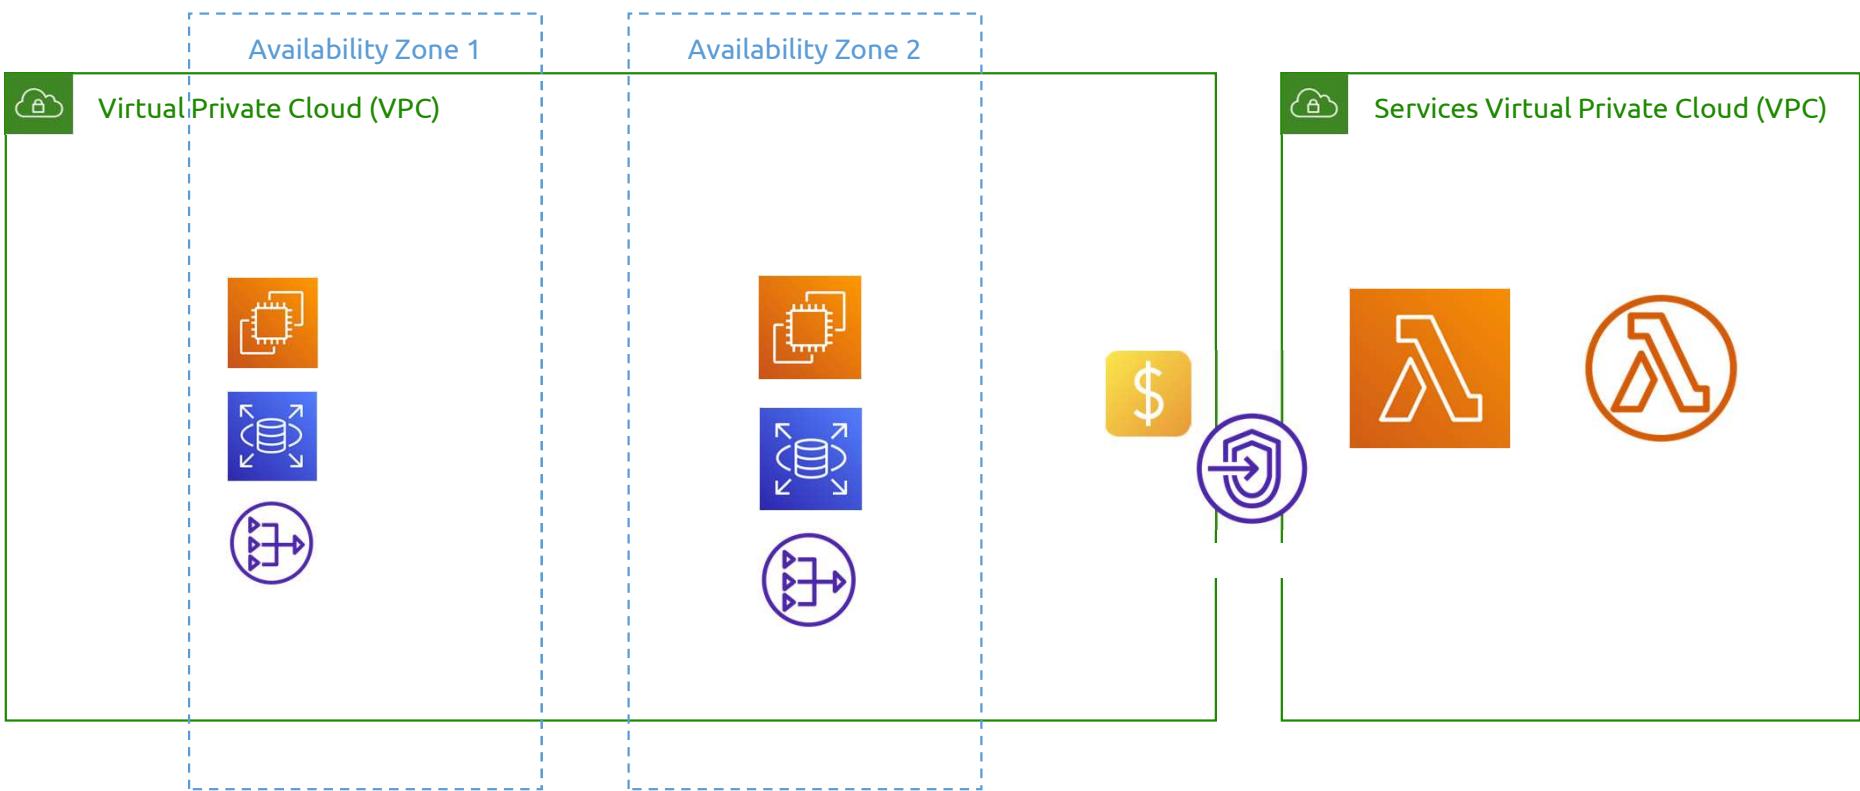
Lambda Networking



Lambda Networking



Lambda Networking





Summary

Networking



- Lambda functions are placed in a special area in AWS called a **services VPC** which are the **isolated networks** in the AWS cloud that holds certain resources.
- The services VPC has access to the internet and other AWS services.
- The services VPC cannot communicate with resources inside private VPCs.
- Two options for adding networking for Lambda: **Running Lambda inside your own private VPC** or **adding a private connection between your VPC and Lambda** using an interface endpoint.
- Running Lambda in your own VPC requires high availability design, a NAT gateway for internet connectivity, and adding endpoints to access other AWS services.
- Running Lambda in your own VPC also affects the number of concurrent Lambda executions and adds to the cost.



KodeKloud

Reserved and Unreserved Concurrency



Concurrency



1



Concurrency

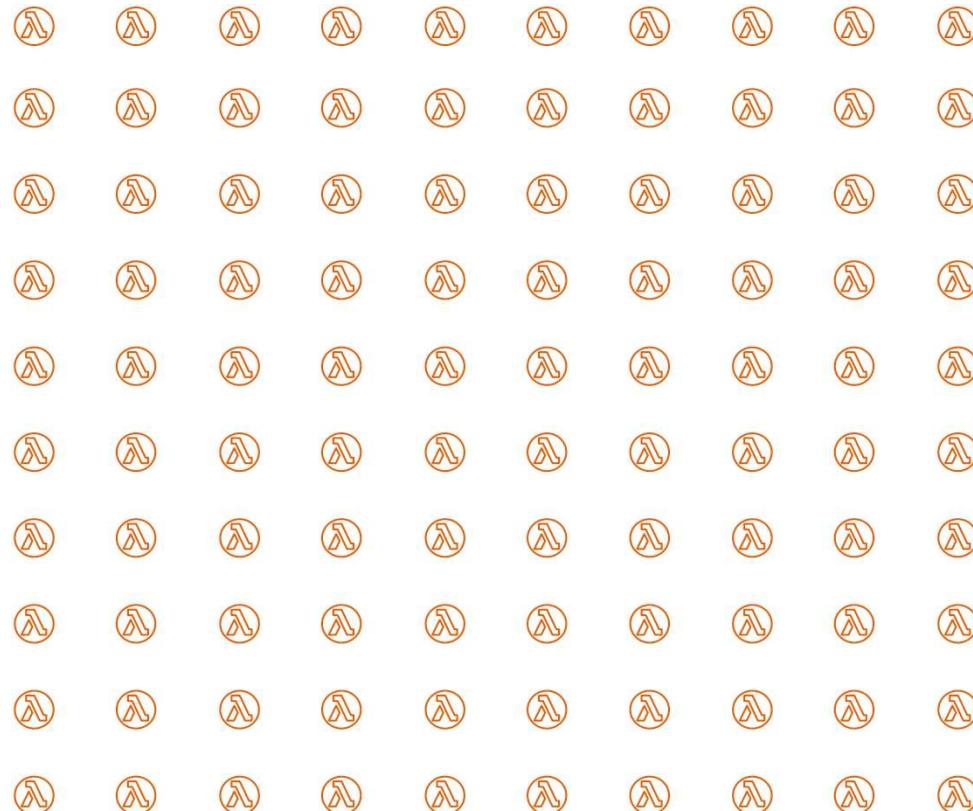


10





Concurrency



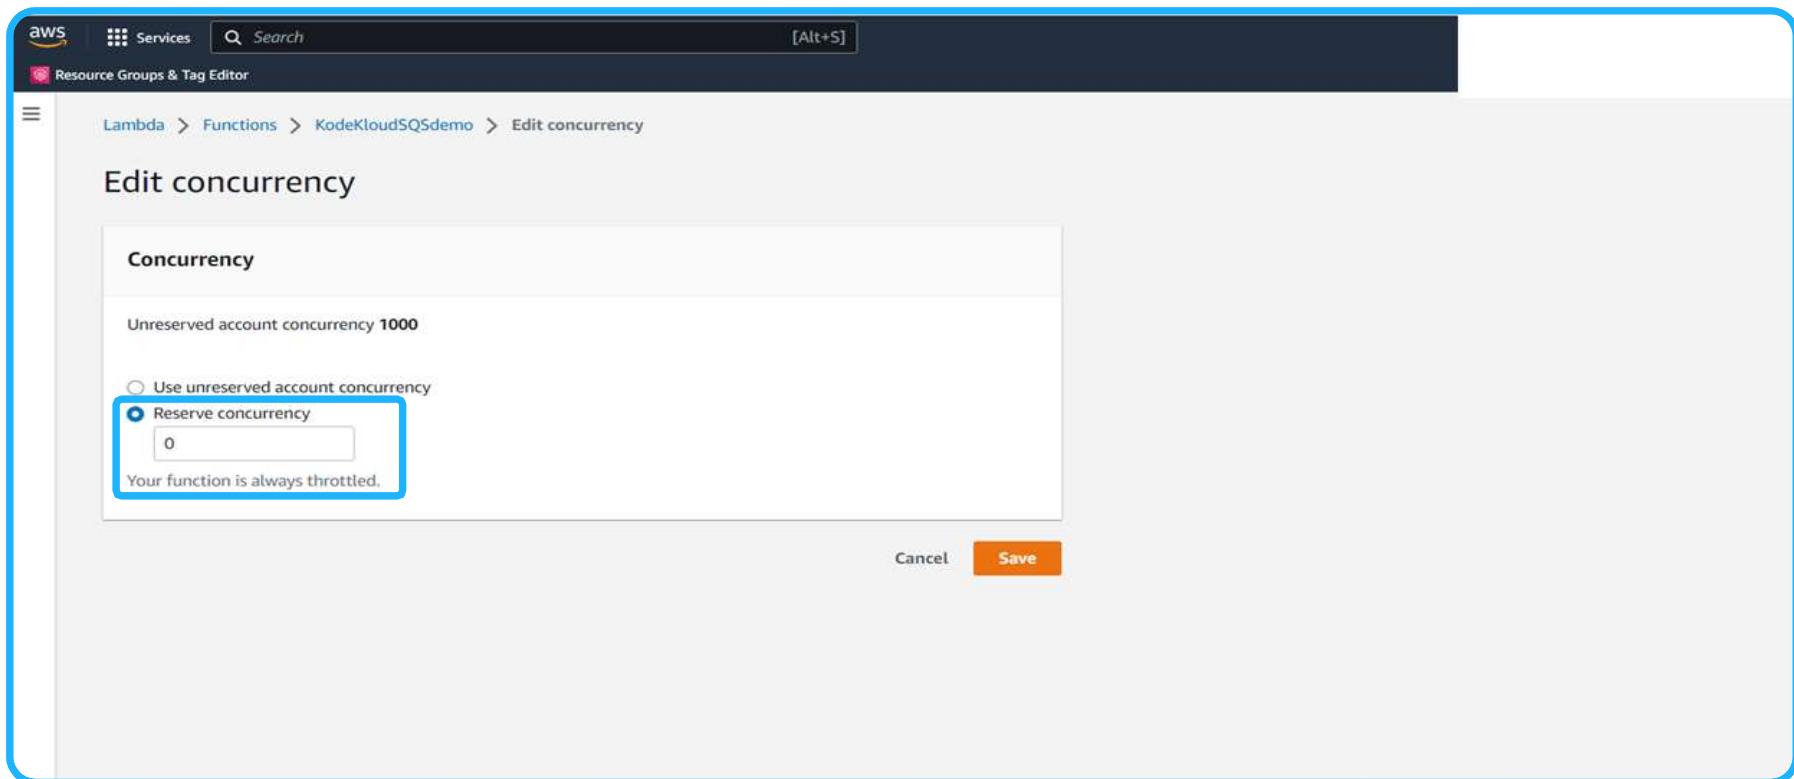
100

Unreserved Concurrency

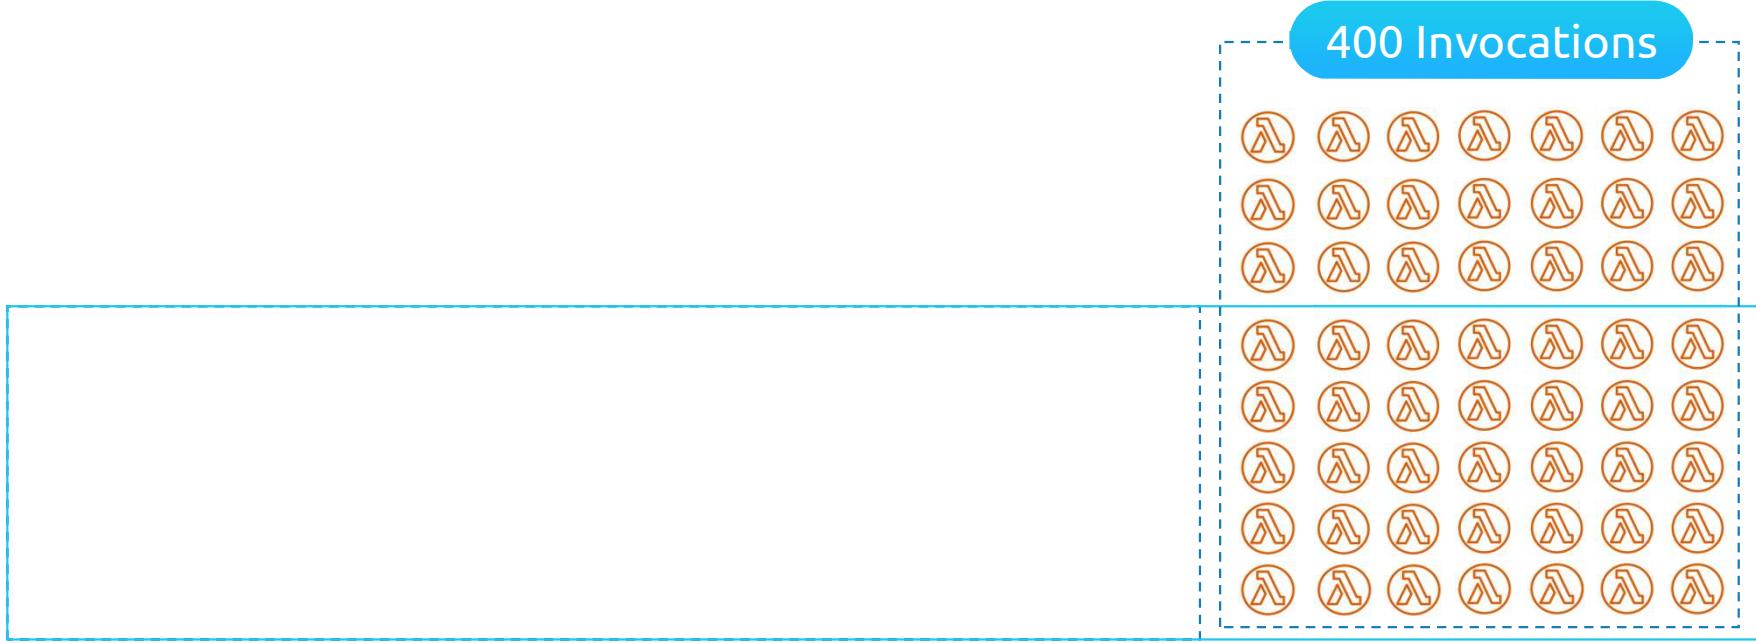
The screenshot shows the AWS Lambda Configuration page. The left sidebar has a 'Concurrency' section selected. The main area displays the 'Concurrency' configuration. It shows 'Function concurrency' set to 'Use unreserved account concurrency' with a value of '1000'. Below this, there's a section for 'Provisioned concurrency configurations' with a note about scaling without latency fluctuations. A table header for 'Provisioned concurrency' includes columns for Qualifier, Type, Provisioned concurrency, Status, and Details. A message indicates 'No configurations' and provides an 'Add configuration' button.

Qualifier	Type	Provisioned concurrency	Status	Details
No configurations				
Add configuration				

Reserved Concurrency



Reserved Concurrency



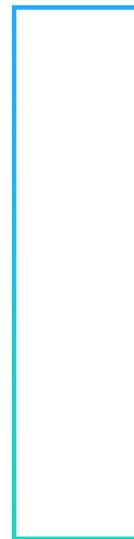
1000 (default account concurrency)

▶ Reserved Concurrency

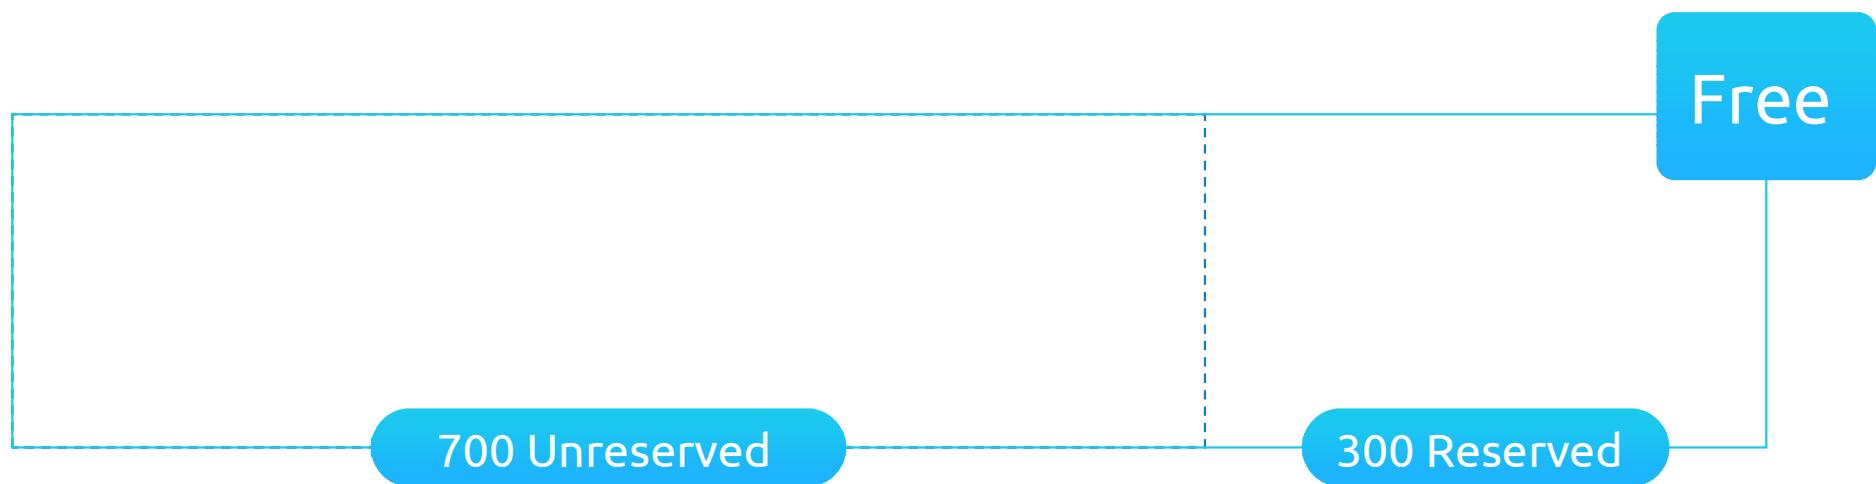
Regulated Concurrency



Overloaded Concurrency



➤ Reserved Concurrency





Reserved Concurrency

1000



Reserved Concurrency

3000

In some Regions

Limited Time



Summary

Concurrency



- AWS Lambda has limits on how many functions can run simultaneously.
- Concurrency settings control which functions get priority within those limits.
- You can configure concurrency settings in the Configuration tab within your function settings.
- [Unreserved Account Concurrency](#) represents the maximum number of Lambda functions you can run simultaneously across your account, usually set to 1000.



Summary

Concurrency



- › Reserved concurrency sets aside a certain number of invocations available to a function at all times but also sets a maximum ceiling for that function.
- › Reserved concurrency guarantees a certain number of invocations available to a function at all times but also sets a maximum ceiling for that function.
- › Setting reserved concurrency is free and can help prevent overloading downstream resources.
- › You can temporarily scale higher than your unreserved concurrency limit, up to 3000 in some regions, for a limited time.



KodeKloud

Provisioned Concurrency

Are there any downsides?

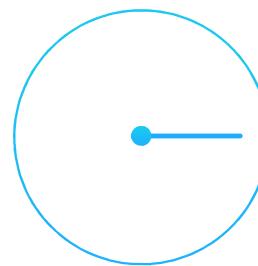
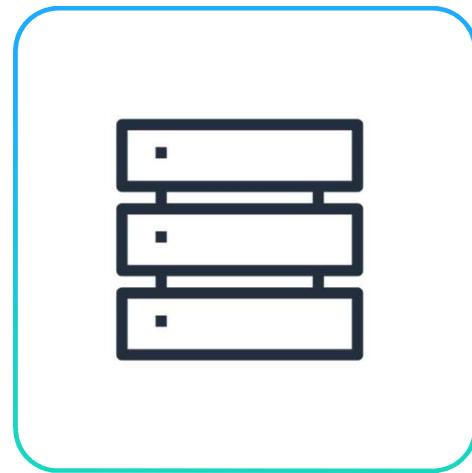
Are we removing benefits from servers?



Lambda Benefits



➤ Physical Server Benefits



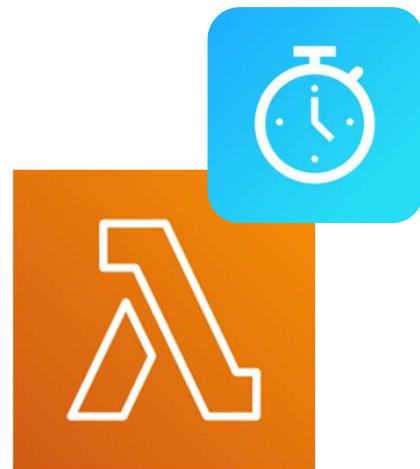


Event Triggers

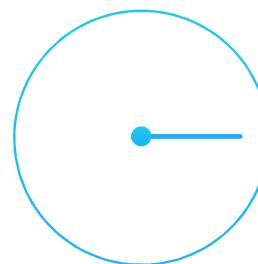
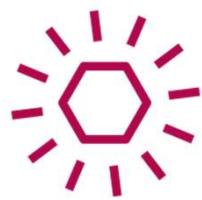




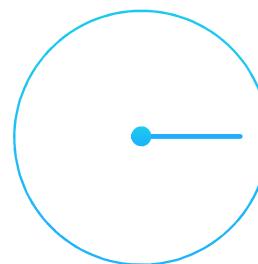
Execution Latency



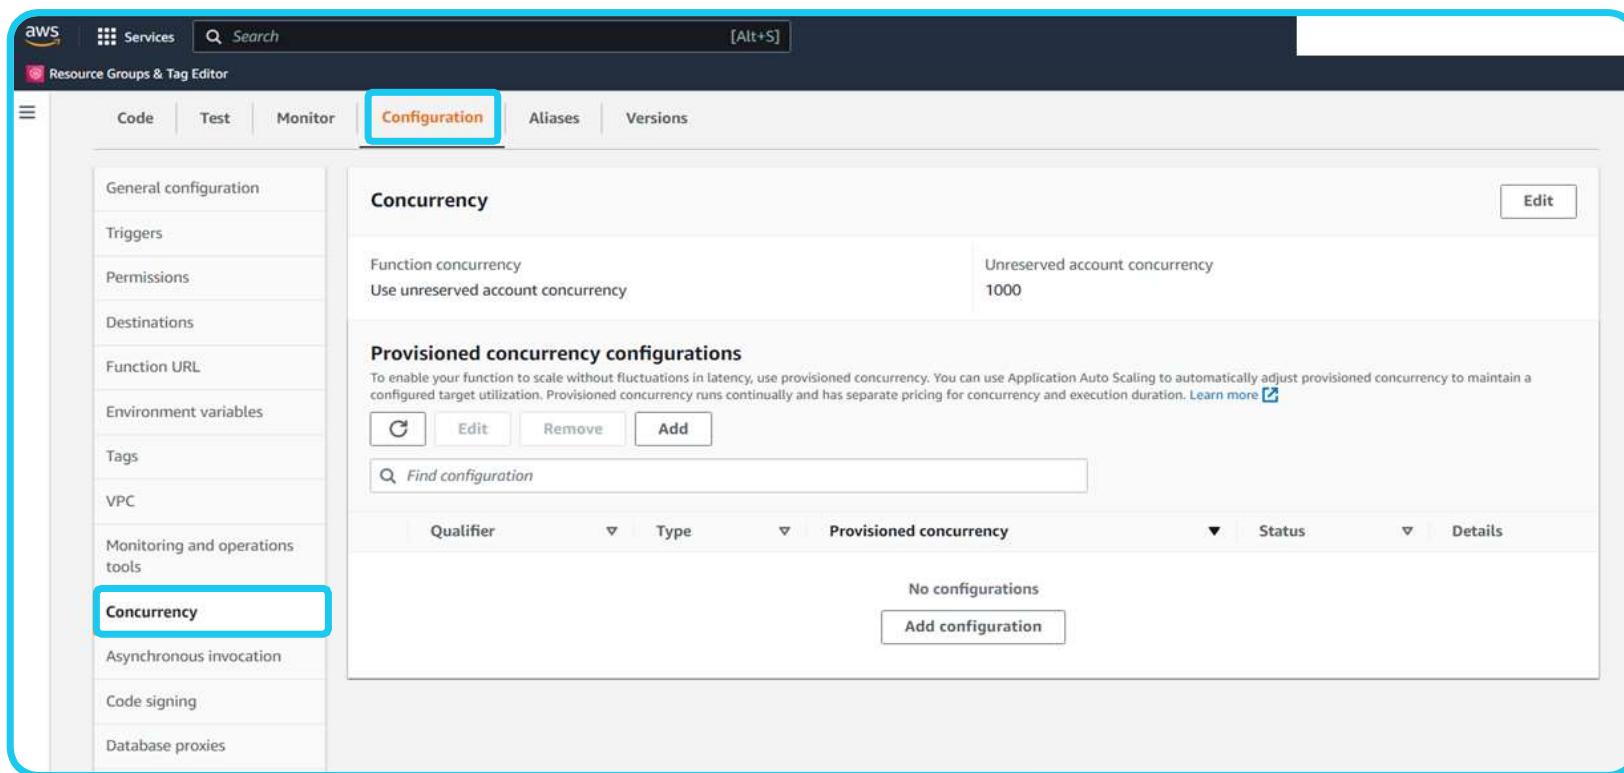
▶ Provisioned Concurrency



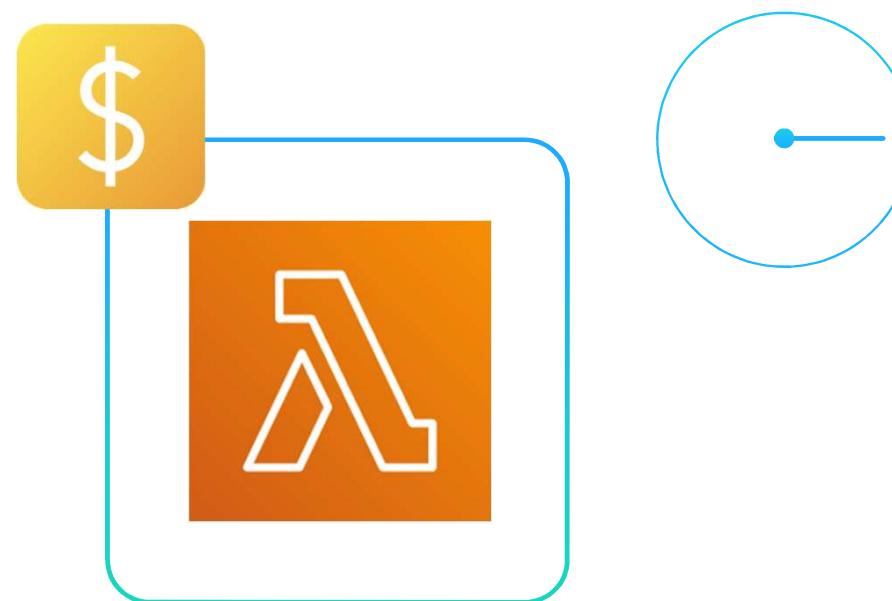
▶ Provisioned Concurrency



Provisioned Concurrency



▶ Provisioned Concurrency





Summary



Provisioned Concurrency

- › Provisioned Concurrency is a solution to avoid execution latency in Lambda, where you tell Lambda to keep a certain number of functions running at all times and ready to execute.
- › Execution latency might occur due to cold starts if the function wasn't running before the event.
- › Provisioned Concurrency can be configured using the concurrency option in the function's configuration tab.
- › Provisioned Concurrency adds cost as it requires functions to be kept running permanently.
- › Aliases and versions are used to assign Provisioned Concurrency.
- › The benefits of removing servers include not having to patch, maintain, and operate them and not having to wait for them to be built to deliver applications.



KodeKloud

Lambda Containers

New

Lambda Containers



Lambda Containers

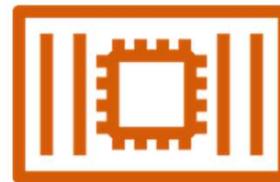




Lambda Containers



?



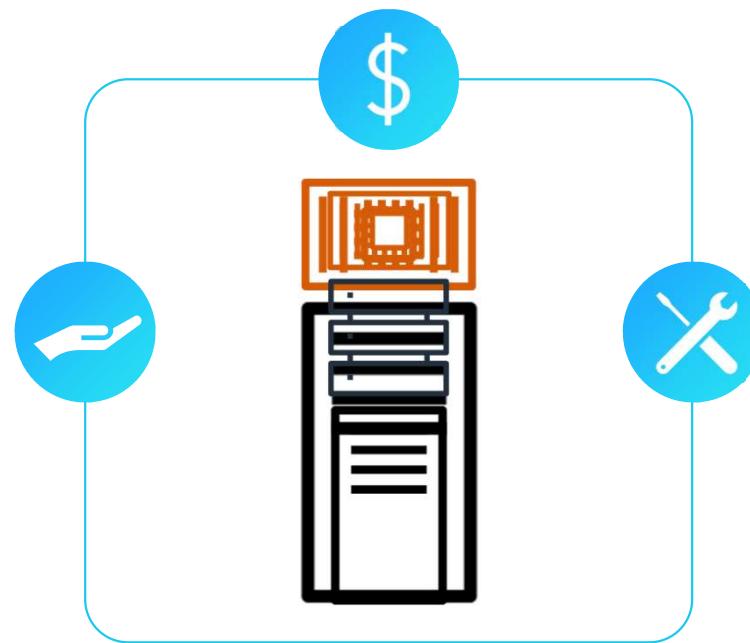


Lambda Containers



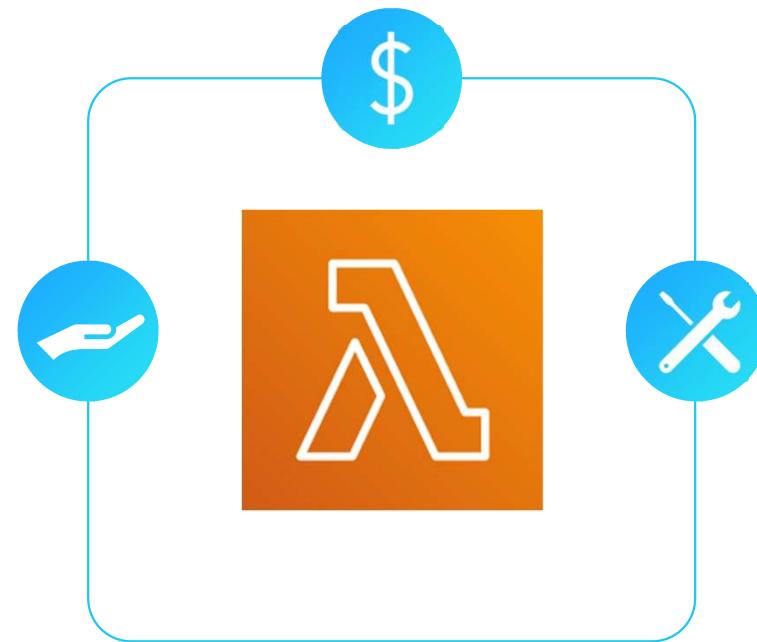


Lambda Containers





Lambda Containers





Lambda Containers



Kubernetes



ECS Autoscaling





Lambda Containers





Lambda Containers





Lambda Containers





Lambda Containers

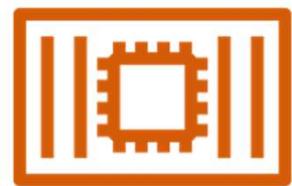


— 10 GB

— 250 MB



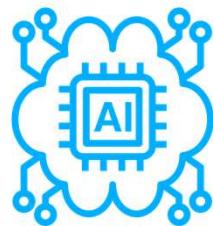
Lambda Containers



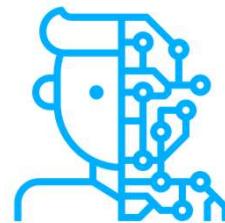
— **10 GB**



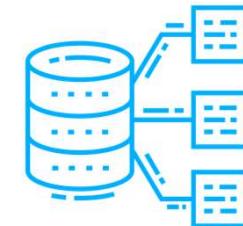
Lambda Containers



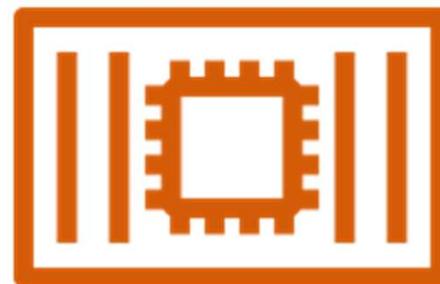
Artificial Intelligence



Machine Learning



Big Data Analytics





Lambda Containers

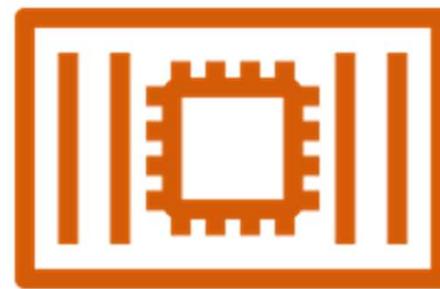
Base image



Code

Base image

Runtime Interface
Emulator



Container App



Summary

Lambda Containers



- › Containers package all **necessary code** and **dependencies** to run a program into a small bundle that can be **run on any machine or OS**.
- › Containers are popular due to their **portability** and convenience but are still typically run on servers.
- › Lambda allows for the running of containers, which offers benefits such as **removing server-related issues**, **automatic scaling**, and **cost efficiency**.
- › Containers on Lambda allow for setting environment variables, specifying startup parameters, and connecting to a VPC.



Summary



Lambda Containers

- AWS Fargate can be used with containers on Lambda for even more **cost efficiency** and **operational simplicity**.
- Lambda supports container file sizes up to **10GB**, allowing powerful use cases such as AI and big data analytics.
- AWS provides base images and a runtime interface emulator to help users integrate containers with Lambda.
- Containers on Lambda may only be suitable for some use cases due to networking limitations or specific dependencies that cannot be packaged.



KodeKloud



Create Basic Function Demo

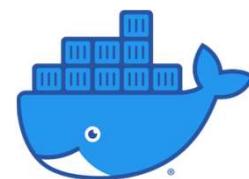
Environment



AWS CLI



Visual Studio Code
(or equivalent)



Docker Desktop
(or equivalent)



AWS Console



Container Demo

Links:

<https://docs.aws.amazon.com/lambda/latest/dg/images-create.html#images-create-from-base>

<https://aws.amazon.com/blogs/aws/new-for-aws-lambda-container-image-support/>



Create a Basic Function Demo

Environment



AWS CLI



Visual Studio Code
(or equivalent)



AWS Console



Create a Basic Function Demo

Links

AWS Command Line Interface Program
<https://aws.amazon.com/cli/>

Lambda CLI Documentation
<https://aws.amazon.com/lambda/latest/dg/gettingstarted-awscli.html>



KodeKloud